

Affine Runtime Calibration for Drift-Adaptive Gottesman–Kitaev–Preskill Decoding

Anonymous authors

Abstract

Approximate Gottesman–Kitaev–Preskill (GKP) error correction converts continuous modular syndrome measurements into displacement corrections. This conversion becomes difficult when finite-energy states, noisy measurement chains and drifting operating conditions move the syndrome distribution over time. Fully learned per-shot decoders can be hard to bound in a real-time control stack, whereas fixed affine correction laws are fast but brittle. We therefore formulate drift-adaptive GKP decoding as affine runtime calibration: the latency-critical path is restricted to applying $\Delta_t = K_t s_t + b_t$, and a slower loop updates (K_t, b_t) from recent analog syndrome statistics.

This formulation targets a gap between two mature lines of work. GKP and bosonic-code decoders increasingly exploit analog reliability information, but the updated object is often an outer-code weight, a message-passing likelihood or a full decoder state. Real-time hardware decoders, in contrast, succeed by making the online datapath small, clockable and reproducible. Our contribution is to connect these requirements at the physical correction layer: recent analog statistics are used to recalibrate a low-dimensional affine displacement surface, while each shot still sees a deterministic matrix-vector rule.

In a predeclared four-scenario simulation protocol, the hybrid residual branch ranks lowest by mean protocol-defined `final_ler_mean` in all four tested scenarios under the two-repeat main table. The scenario-wise mean reduction relative to the unscented Kalman filter is 1.8–2.9%, with a 2.3% average across the tested scenarios. Two completed scenario rows, `static_bias_theta` and `linear_ramp`, provide 12-pair positive interval checks: paired- t lower bounds of 0.014778 and 0.021215, and bootstrap 95% lower bounds of 0.014933 and 0.021405. Pooled and all-scenario interval claims remain unsupported. Ablations and a six-seed mechanism analysis indicate that histogram dynamics and teacher-side features affect the result, but they do not prove a unique causal mechanism or teacher necessity. A supplementary statistical-calibration analysis further shows that a non-neural estimator can populate the same affine interface, supporting the contract rather than a specific neural module. Controlled nearest-syndrome, oracle-affine and wrapped-Gaussian checks separate hard-correction risk, estimator error and local affine-model limits. Additional diagnostics connect the residual-boundary event to a Pauli-style surrogate channel, show unchanged boundary-crossing rates under Q4.20 fixed-point emulation in the controlled sample set, and reduce the fast path to four multiplications, four additions and six stored state scalars per shot. Board timing, resource use, power and source-vs-board measurements are not evaluated here and require direct hardware measurements. The results therefore support a systems-method claim: low-dimensional affine calibration is a plausible organizing principle for drift-adaptive GKP fast-path decoding, while deployment-level and broad-benchmark claims require direct device and repeat-expanded validation.

1 Introduction

Bosonic quantum error correction encodes information in oscillator degrees of freedom and can exploit continuous measurement records that are unavailable in binary stabilizer readout. The GKP

code is a prominent example: shift errors are inferred from modular quadrature syndromes and corrected by displacement operations [1, 2]. In an idealized setting, the decoder can treat each measured syndrome as a local displacement estimate. In a physical setting, however, approximate finite-energy GKP states, noisy auxiliary measurements and hardware calibration drift alter the distribution seen by the decoder over time.

The resulting bottleneck is not only statistical accuracy, but the interface between estimation and real-time control. A drift-aware decoder should exploit recent analog syndrome information, yet a decoder intended for a hardware control stack must expose a small, testable and latency-bounded correction surface. End-to-end neural decoding can be expressive, but it obscures worst-case timing, fixed-point degradation and rollback behaviour. A fixed affine correction law has the opposite profile: it is fast and interpretable, but it cannot by itself follow bias, variance or orientation drift in the effective GKP noise state.

We separate these two responsibilities. The fast loop is limited to a two-quadrature affine rule,

$$\Delta_t = K_t s_t + b_t, \quad (1)$$

where s_t is the measured syndrome and (K_t, b_t) are the active runtime parameters. The slow loop updates these parameters from a recent syndrome histogram, teacher estimate or statistical calibration rule. The learned component is therefore not the per-shot decoder; it is one possible slow-loop module for recalibrating a low-dimensional physical displacement rule. Different estimators can be compared only after they commit the same type of (K, b) parameters to the deterministic fast path.

We test whether a slow calibration module can improve a deterministic affine GKP fast path under controlled drift while preserving a small, verifiable runtime interface. Under a predeclared controlled simulation protocol, the main-table hybrid residual branch lowers the protocol-defined `final_ler_mean` in every tested drift scenario relative to the strongest baseline in this protocol, UKF. The present validation is simulation-based and leaves board-level timing, resource use and source-vs-board agreement as explicit hardware-stage measurements.

The comparisons in this paper are therefore tiered. The direct numerical comparison is the protocol-defined `final_ler_mean` ranking within the controlled drift benchmark. The fidelity discussion is a residual-boundary surrogate that explains how half-lattice crossings relate to channel language, not a finite-energy logical-channel reconstruction. The latency and resource discussion is a datapath argument based on operation count, fixed-point software parity and explicit hardware measurement requirements. Maintaining this separation lets the reader distinguish the current simulation result from the channel-level and device-level standards used in adjacent work.

This tiering also defines the specific advantage being tested. The paper does not argue that an affine rule is a universal GKP decoder, nor that a small CNN should replace established soft-information or learned decoders. It argues that a slow calibration loop can absorb nonstationary syndrome statistics and commit a compact (K, b) surface whose online correction is designed to be easier to bound, quantize and log than a per-shot neural or multi-branch posterior calculation; hardware reproduction remains a future measurement target. The main evidence therefore concerns adaptation under drift and contract simplicity, not end-to-end fault-tolerance thresholds.

The new element is the interface between GKP analog statistics and a bounded runtime surface. First, the updated object is the physical-layer affine correction (K, b) , not an outer-code graph weight or a full neural decision rule. Second, the slow-loop estimator can be neural, statistical or teacher anchored, provided that it commits the same low-dimensional surface to the fast path. Third, the evidence is organized around that shared contract: the main simulation ranking tests whether residual calibration improves `final_ler_mean`, while the cost, fixed-point and traceability analyses test whether the resulting fast path is a plausible hardware verification target.

Runtime-consistent affine calibration

Slow-loop estimation updates low-dimensional parameters; the per-shot path remains deterministic.




Figure 1: Overview of runtime-consistent affine calibration. The slow calibration loop updates low-dimensional (K, b) parameters from recent syndrome statistics, while the latency-critical fast path remains the deterministic affine operation in Eq. (1). The dashed hardware-facing surface marks future timing, fixed-point and resource measurements, not current hardware results.

The main contributions are:

1. A runtime-consistent formulation of drift-adaptive GKP decoding as affine calibration of (K, b) , with the per-shot correction constrained to Eq. (1).
2. A residual branch, instantiated here with teacher anchoring in the main comparison, that predicts a bounded bias calibration term while leaving the fast path deterministic.
3. A predeclared four-scenario controlled simulation evaluation showing that the hybrid residual branch ranks lowest by mean `final_ler_mean` in all tested scenarios under the current two-repeat protocol, plus two scenario-level 12-paired-repeat interval checks that strengthen the `static_bias_theta` and `linear_ramp` rows without yet establishing an all-scenario interval result.
4. A metric ladder separating protocol-defined `final_ler_mean`, residual-boundary channel language, analytical fast-path cost, fixed-point software parity and the hardware measurements needed before deployment-level claims.
5. A comparison framework that positions the method against analog GKP decoding, adaptive-prior calibration, learned QEC modules and real-time FPGA decoders without treating incompatible metrics as a single leaderboard.

2 Related Work

GKP analog information. Analog GKP syndrome values carry reliability information that can improve decoding beyond hard thresholding. Prior work has used analog information in high-threshold GKP error correction, surface-GKP decoding and bosonic-QLDPC-GKP settings [3, 4, 5, 6, 7, 8]. The present work does not claim novelty in using analog information. Its distinction is the layer at which the information is used: recent analog syndrome statistics update a physical-layer affine displacement rule, rather than only an outer-code matching weight or message. In other words, analog information is the established input class; the new design question is where that information is committed in a hardware-facing decoder stack. The histogram window used here is also not a Bayesian multi-round approximate GKP memory decoder and does not model noisy auxiliary-state correlations [9, 10]. It is an effective-noise calibration signal for a committed physical affine surface.

Adaptive priors and calibration-aware decoding. Decoder performance depends on the match between assumed and actual noise. Adaptive weight estimation, syndrome-statistics estimation and calibrated decoders address this mismatch by updating priors or graph weights from observed data [11, 12, 13, 14, 15]. Our formulation follows the same motivation, but constrains the updated object to the GKP affine fast-path parameters and commits those parameters through a stage-and-commit runtime contract. The calibrated quantity is therefore not a decoder score or graph edge alone, but the (K, b) surface consumed by the physical correction rule.

Learned low-latency QEC modules. Machine-learning-assisted decoders and pre-decoders can reduce latency or improve accuracy when their role in a classical decoding stack is explicit [16, 17, 18, 19, 20, 21]. Structured neural decoders, soft-information decoders and calibration-driven control updates have also been studied in bosonic, surface-code and hardware-data settings [19, 20, 21, 22]. The residual module studied here is narrower than a full neural decoder: it predicts a bounded calibration term for a low-dimensional affine surface and is kept outside the per-shot critical path. The novelty claim is therefore not “neural GKP decoding”, but the restriction of adaptation to a committed affine fast-path surface. The closest learned-decoder contrasts are calibration-conditioned or real-time neural modules, such as folded FiLM-style calibration and FPGA neural decoding [18, 23]. Those systems place learned computation closer to the decoder or feedback path; here, learning is confined to slow-loop calibration and the fast path remains an affine displacement rule.

Real-time and FPGA QEC decoders. Hardware QEC decoders for stabilizer-code settings include lookup-table, union-find, matching-accelerator, clustering, greedy and FPGA-tailored message-passing systems [24, 25, 26, 27, 28, 23, 29, 30]. These works define the standard for deployment claims: closed-loop latency, resource use, timing closure and hardware integration. This paper addresses a different measurement layer: it treats the GKP affine fast-path contract as the object to be isolated in controlled simulation and fixed-point software checks, with hardware-facing metrics reserved for board-level implementation. This distinction also fixes the comparison unit. The prior FPGA systems evaluate complete stabilizer-code decoding workloads and report device-level timing, memory or integration evidence. The present study evaluates a physical-layer correction surface that would sit upstream of such a broader decoder stack. Its hardware relevance is therefore the compactness and observability of the online rule: the affine update is small enough to be fixed-point checked, logged and compared against a later source-vs-board test target. Hardware systems such

Topic	Prior scope	Distinction of this work
GKP analog decoding	Uses continuous syndrome values or reliability weights for GKP or outer-code decoding.	Uses recent analog histograms to calibrate the physical-layer affine displacement rule itself.
Adaptive priors	Updates noise weights, graph weights or decoder priors from syndrome statistics.	Maps the estimated state to committed (K, b) parameters consumed by a deterministic fast path.
Learned QEC modules	Places neural modules before, inside or after a classical decoder.	Keeps the neural branch in the slow calibration loop rather than on the per-shot correction path.
FPGA/real-time decoders	Reports deployed low-latency or hardware-integrated decoders in stabilizer-code settings.	Provides a software-tested GKP physical-layer calibration contract; hardware results require future board-level measurements.

Table 1: Topic-level positioning. The contribution is the runtime-bounded affine calibration contract for a drifting GKP physical correction layer, not a dismissal of prior analog, learned or hardware-decoder work.

as lookup-table, union-find, clustering and neural-FPGA decoders report latency, resource, memory or closed-loop evidence as device measurements [24, 25, 26, 28, 23]; this manuscript instead defines the affine target that such a measurement would need to validate.

The closest prior-work cluster is therefore not a single decoder family, but the intersection of three constraints: analog GKP reliability information, adaptive noise or prior calibration, and real-time decoder implementation. Existing GKP soft-information work establishes the value of continuous syndrome data; adaptive-prior and calibration-conditioned decoders establish that decoder parameters should track drift; and FPGA/QEC runtime work establishes that deployment claims require bounded datapaths and measured latency. The present work is positioned at their intersection: it updates a physical-layer affine correction surface while keeping the per-shot correction path deterministic.

The comparison is metric-oriented rather than leaderboard-style. Prior work reports different code families, validation layers and endpoint metrics, so Table 2 identifies the measurement standard each line establishes and the narrower metric reported here; the expanded literature crosswalk is provided in Tables 34 and 35 in the appendix. We use three rules when reading these external numbers. First, endpoint metrics are not exchanged across layers: surface-GKP thresholds, logical-channel infidelities, real-processor LER and FPGA latency are standards for comparison, not measurements produced by this study. Second, the adaptive object must be named: prior work often updates an outer-code graph, a decoder prior, a neural decoder state or a message-passing likelihood, whereas this manuscript updates the physical-layer affine surface (K, b) . Third, the runtime surface must be separated from the estimator: the present fast path remains a deterministic clipped affine rule even when the slow-loop estimator is learned.

Taken together, the comparison defines a narrow but concrete advantage claim. The manuscript does not compete with surface-GKP or QLDPC-GKP decoders on end-to-end overhead, and it does not compete with real-time FPGA decoders on measured device latency. Instead, it tests whether drift adaptation can be expressed as a small physical-layer affine update whose online arithmetic is much simpler than posterior branch search while still improving the protocol-defined

LER proxy under controlled drift. The closest adjacent neural-FPGA demonstration is therefore a standard to validate against, not a baseline that this study has already matched. The design trade-off tested here is architectural and decomposes into three parts: the correction surface adapts under controlled drift, the latency-critical operation remains a clipped affine matrix-vector rule, and the implementation target is small enough to be checked against source-level vectors. This gives the manuscript a route to three later hardware-facing metrics: lower analytical per-shot arithmetic count than the branch-aware posterior references evaluated in the controlled cost model, a smaller source-level verification target than a per-shot neural decoder and a clear source-vs-board agreement target for an FPGA implementation.

This protocol also explains which comparisons are most probative for the present claim. Analog-GKP and calibration-aware QEC work establish that continuous syndrome information and drifting calibration states matter. They do not by themselves establish a hardware-friendly physical correction interface. Learned-decoder and real-time-FPGA work establish the evidence bar for LER and latency claims. They do not remove the need for a small, verifiable fast path when the goal is a GKP correction module that can later be implemented with fixed-point arithmetic. The contribution is therefore not a new universal decoder, but a way to expose drift adaptation through a low-dimensional surface whose online arithmetic, state and future board check are explicitly bounded.

For clarity, the manuscript uses a metric ladder rather than a single all-purpose performance number. The primary result is the protocol-defined `final_ler_mean` ranking and paired descriptive delta. The residual-boundary surrogate is only a bridge from half-lattice crossing events to channel language, not a finite-energy logical-channel fidelity. The fast-path latency argument is currently an analytical and software-emulation argument, not an FPGA timing result. Table 3 separates the metrics reported in this study from the metrics that define the standard for later validation.

This ladder is the intended reading guide for all cross-paper comparisons in the manuscript. The three quantitative axes should be read as a triangulation, not as interchangeable endpoints. The `final_ler_mean` rows test whether adaptive calibration changes the protocol-defined boundary-crossing rate under drift. The residual-boundary surrogate then asks whether the same q/p events can be reported in channel language without pretending to reconstruct a finite-energy logical channel. The operation-count and fixed-point rows ask whether the resulting correction rule has the size and numerical stability expected of a future real-time datapath. A positive result on only one axis would be weak: lower `final_ler_mean` without a channel-language bridge would be hard to compare with GKP fidelity work, a surrogate fidelity without the `final_ler_mean` protocol would be detached from the reported benchmark, and a low arithmetic count without parity checks would not justify a hardware-facing claim.

3 Problem Setting

3.1 Approximate GKP syndrome model

The ideal square-lattice GKP code represents logical states as combs in phase space. In the position quadrature, $|0_L\rangle$ and $|1_L\rangle$ occupy alternating lattice points separated by $\sqrt{\pi}$, while stabilizer periods are $2\sqrt{\pi}$. Approximate GKP states replace the Dirac comb by finitely squeezed peaks with a finite-energy envelope, so both state preparation and readout add continuous noise to the modular syndrome. For example, one may view an approximate position-basis codeword schematically as

$$|\tilde{0}_L\rangle_q \propto \sum_{n \in \mathbb{Z}} \exp\left[-\frac{(2n)^2\pi}{2\Delta_{\text{env}}^2}\right] \int dq \exp\left[-\frac{(q - 2n\sqrt{\pi})^2}{2\Delta_{\text{pk}}^2}\right] |q\rangle, \quad (2)$$

with the odd lattice sites representing $|\tilde{1}_L\rangle_q$. The peak width Δ_{pk} represents finite squeezing and the broad envelope Δ_{env} limits oscillator energy. The same description applies in the conjugate quadrature after the corresponding Fourier rotation. The exact envelope, normalization and circuit depend on the physical preparation route, but the decoder-facing consequence is common: finite squeezing turns a sharp lattice-cell decision into a tail probability near the half-lattice boundary. A small change in residual scale or bias can therefore produce a disproportionately large change in boundary-crossing probability when the residual distribution sits near that decision surface.

This tail sensitivity is the link between approximate-state physics and the software metric used below. In the usual Gaussian peak approximation, modular squeezing can be summarized by a width parameter such as $S = -10 \log_{10} \Delta_{\text{pk}}^2$, while the finite-energy envelope controls how much probability leaks into distant lattice peaks. After a syndrome measurement and correction, a residual r_j^+ that remains well inside the local cell is still compatible with the intended logical branch; a residual that crosses $\pm\lambda/2$ changes the nearest branch assignment. For a zero-mean Gaussian residual with variance σ_r^2 , the corresponding one-quadrature crossing probability is approximately $2Q(\lambda/(2\sigma_r))$, where Q is the Gaussian upper-tail function, so modest changes in residual scale become exponentially visible near the boundary. This is why a drift-adaptive affine surface can matter even when its correction rule is simple: it shifts the post-correction residual distribution away from the cell edge rather than attempting to reconstruct the full finite-energy wavefunction.

This finite-energy picture also fixes the level of abstraction used in the paper. The ideal code treats displacements separated by a stabilizer period as equivalent, whereas an approximate code supplies only a noisy, periodic measurement of the displacement. The decoder therefore does not observe the absolute shift e_t . It observes a modular residual together with preparation, ancilla and readout noise, then chooses a correction that is safe only within the selected lattice cell. In this manuscript, a q or p event is counted when the post-correction residual crosses the half-period boundary of that normalized cell. This is the operational origin of the reported `final_ler_mean` proxy and of the later Pauli-style surrogate tables.

The syndrome readout is therefore modular rather than absolute. If a pre-correction displacement e_j occurs in quadrature $j \in \{q, p\}$, the readout exposes e_j only modulo the lattice period and with additive preparation and measurement uncertainty. The correction is successful within the modeled cell when the post-correction residual $r_{j,t}^+$ remains inside the interval $[-\lambda/2, \lambda/2)$. The manuscript's per-round event indicator can be written as

$$\ell_t = \mathbb{1}\{|r_{q,t}^+| \geq \lambda/2 \text{ or } |r_{p,t}^+| \geq \lambda/2\}, \quad (3)$$

and `final_ler_mean` is the empirical mean of this indicator under the stated software protocol. This definition is intentionally narrower than a code-distance logical failure probability: it measures whether the physical fast-path correction leaves the residual in the wrong normalized cell, not whether an outer code or a full finite-energy logical channel has failed.

To avoid mixing oscillator and software coordinates, the remainder of the paper uses one normalized residual convention. The conventional square-lattice GKP description above explains the physical code. The runtime equations below use the software residual coordinate in which the wrap period is $\lambda = \sqrt{2\pi}$. All q/p crossings, clipping thresholds and source tables are expressed in that normalized coordinate.

Let the accumulated two-quadrature displacement before correction be $e_t = (e_{q,t}, e_{p,t})^\top$. The implementation uses the lattice period $\lambda = \sqrt{2\pi}$ for syndrome wrapping and half-lattice residual boundaries. A compact measurement model is

$$\tilde{s}_t = \text{wrap}_{[-\lambda/2, \lambda/2)}(e_t) + \eta_t^{\text{meas}} + \eta_t^{\text{GKP}}. \quad (4)$$

The two noise terms summarize effective readout and finite-squeezing contributions in the software model. They are not fitted to a calibrated finite-energy GKP device channel in this study. The constant λ should be read as the coordinate convention used by the software residual model, not as a new physical code definition. Standard GKP presentations often describe logical separation and stabilizer periods in the oscillator quadrature normalization natural to the preparation circuit. Here all measured syndromes, accumulated residuals and logical-boundary tests are expressed in a normalized two-quadrature residual coordinate with period $\sqrt{2\pi}$. This convention makes the runtime rule and the supporting data unambiguous: every wrap, clipping threshold and boundary-crossing event in the manuscript is evaluated against the same λ , while squeezing in dB, photon number and channel-specific finite-energy parameters remain outside the present calibrated model. Equivalently, the manuscript works after a fixed coordinate rescaling from the usual oscillator-quadrature convention into the software residual coordinate; all reported q/p crossings are crossings of this normalized cell, not separate claims about a device-specific physical quadrature calibration. The wrap operation is the essential ambiguity: the same reported syndrome can correspond to different lattice branches. If the posterior mass is concentrated near one branch, a local affine estimator is a reasonable low-latency approximation. For jointly Gaussian local variables e and s , the linear-MMSE estimate is

$$\hat{e} = \mu_e + \Sigma_{es}\Sigma_{ss}^{-1}(s - \mu_s) = Ks + b, \quad K = \Sigma_{es}\Sigma_{ss}^{-1}, \quad b = \mu_e - K\mu_s. \quad (5)$$

This derivation motivates the affine fast path while also defining its limits. More explicitly, fix a lattice branch $m_t \in \mathbb{Z}^2$ and define the unwrapped local residual $r_t = e_t - \lambda m_t$. Conditional on the effective noise state θ_t^{noise} and on this branch assignment, suppose the local pair (r_t, s_t) is approximated by a joint Gaussian with moments $(\mu_r, \mu_s, \Sigma_{rs}, \Sigma_{ss})$. The minimum-mean-square affine correction is then

$$\hat{r}_t = \mathbb{E}[r_t \mid s_t, \theta_t^{\text{noise}}, m_t] \approx \mu_r + \Sigma_{rs}\Sigma_{ss}^{-1}(s_t - \mu_s) = K(\theta_t)s_t + b(\theta_t). \quad (6)$$

The affine rule used by the fast path is therefore not introduced as a universal wrapped decoder. It is the single-branch local conditional-mean surface that remains after the slow loop has compressed recent syndrome statistics into a small committed parameter bank. The scientific question in this paper is whether such a surface can remain useful under controlled drift while preserving a deterministic low-latency datapath. With the sign convention used in the emulator, the affine output is the correction displacement applied to the pre-correction shift. The post-correction local residual is

$$r_t^+ = e_t - \Delta_t - \lambda m_t, \quad \Delta_t = K_t s_t + b_t, \quad (7)$$

where m_t is the selected lattice branch. Boundary events in the metric section are defined on r_t^+ , so the affine surface is judged by the residual it leaves after correction rather than by parameter error alone.

Proposition 1 (Local affine fast path). *Fix an effective noise state θ and a lattice branch m . Let (r, s) denote the corresponding unwrapped local residual and wrapped syndrome, with finite second moments and nonsingular Σ_{ss} . Among all affine corrections $a(s) = As + c$, the unique minimum-mean-square correction is*

$$A^* = \Sigma_{rs}\Sigma_{ss}^{-1}, \quad c^* = \mu_r - A^*\mu_s.$$

If (r, s) is jointly Gaussian within this conditioned branch, this affine rule is also the conditional mean $\mathbb{E}[r \mid s, \theta, m]$.

Proof. For an affine correction $a(s) = As + c$, the mean-square risk is $\mathbb{E}\|r - As - c\|_2^2$. Differentiating with respect to c gives $c = \mu_r - A\mu_s$. Substituting the centred variables $\bar{r} = r - \mu_r$ and $\bar{s} = s - \mu_s$, the remaining risk is $\mathbb{E}\|\bar{r} - A\bar{s}\|_2^2$. The normal equations are $A\Sigma_{ss} = \Sigma_{rs}$, which yields $A^* = \Sigma_{rs}\Sigma_{ss}^{-1}$ when Σ_{ss} is nonsingular. For a jointly Gaussian pair, the conditional mean is linear with the same coefficient and intercept. $\square$

The proposition is intentionally local. It justifies the committed (K, b) surface as the optimal affine correction after a branch and an effective noise state have been compressed by the slow loop. It does not claim global optimality for wrapped, multi-branch posteriors, nor does it replace a finite-energy logical-channel simulation. Its role is to make the runtime contract mathematically explicit: a slow estimator may update the moments or their learned proxy, while the fast path only evaluates the affine map fixed by the proposition.

This problem setting is deliberately weaker than a full logical-channel simulation. A complete approximate-GKP analysis would track finite-energy envelopes, auxiliary-state noise, channel-specific loss or amplification, and the induced logical Pauli channel. The present model instead treats these effects through an effective wrapped displacement process and asks whether a runtime-constrained correction surface can adapt to controlled nonstationarity. That restriction makes the FPGA-facing question concrete: the fast path receives a wrapped syndrome and a committed affine parameter bank, not a full posterior over lattice branches.

The separation is important for comparing with fidelity-oriented GKP work. Channel-level studies ask how a physical loss, amplification or finite-energy model propagates to logical probabilities or process infidelity. This manuscript asks a narrower systems question: if the physical correction layer is restricted to a committed affine rule, can recent syndrome statistics keep that rule aligned with a drifting effective residual distribution? The answer is reported with a half-lattice boundary proxy and diagnostic channel-language surrogates. It should not be interpreted as logical-channel fidelity because the supporting data do not reconstruct the finite-energy output state or a complete logical process map.

The validity regime can be stated in operational terms. Let B_t be the event that the wrapped syndrome is assigned to the wrong lattice branch after conditioning on the recent histogram, and let c be the slow-loop update window. The affine path is expected to be useful when $\Pr(B_t \mid H_{t-c+1:t})$ is small, the local covariance of the conditioned residual is well described by a single mode, and the drift timescale is long enough that (K_t, b_t) remains informative between commits. It is expected to fail when posterior mass splits across adjacent lattice branches, when finite-energy tails place substantial probability near half-lattice boundaries, or when drift changes faster than the histogram window and commit cadence can track. These three failure modes motivate the paper’s bounded evidence ladder: residual-MSE checks for local affine fit, boundary-crossing `final_ler_mean` for branch-cell errors, holdout drift stress tests for slow-loop lag and a hardware-facing datapath contract for future device validation. The controlled oracle and wrapped-Gaussian analyses in the Results section test this local validity regime rather than replacing a full wrapped decoder benchmark.

3.2 Effective drift state

The slow loop estimates a compact effective noise state rather than a full circuit-level model:

$$\theta_t^{\text{noise}} = (\sigma_t, \mu_{q,t}, \mu_{p,t}, \vartheta_t). \quad (8)$$

Here σ_t summarizes displacement scale, $\mu_{q,t}$ and $\mu_{p,t}$ represent mean bias, and ϑ_t captures covariance orientation. The current benchmark uses four hand-designed drift families: static bias with rotation, linear ramp, step change in variance/orientation and periodic drift. These scenarios test distinct adaptation demands but do not exhaust possible nonstationary GKP noise.

3.3 Evaluation target

The primary metric is `final_ler_mean`, the source-field name for a q/p boundary-crossing count per fast-loop round. Lower values are better. Because the main benchmark data do not include confidence intervals or a formal predeclared holdout set, the main result is interpreted as a predeclared-protocol ranking, not as a broad distributional generalization claim. Controlled holdout stress tests are reported separately in the Results section, and repeat-expanded inferential confidence intervals or predeclared bootstrap intervals for the main rows remain required before stronger statistical language is appropriate.

Metric definition. For one controlled simulation run, `final_ler` is the final value of the cumulative q/p residual-boundary crossing proxy reported by the fast-loop emulator,

$$\text{final_ler} = \frac{N_X + N_Z}{T},$$

where N_X and N_Z are the numbers of q- and p-quadrature residual boundary crossings counted by `LogicalErrorTracker`, and T is the number of fast-loop rounds in the run. The table field `final_ler_mean` is the mean of this value across the available repeats for each scenario and mode; the associated `final_ler_sd` field is a descriptive standard deviation. Because q and p crossings are counted separately, simultaneous crossings contribute two events; the field is therefore an event-count proxy, not a probability bounded by one. This definition is a controlled simulation protocol metric. It is not a hardware logical-error measurement, a logical-channel estimate, a confidence interval or a significance test.

The proxy follows the standard GKP intuition that a residual displacement past the half-lattice boundary corresponds to a logical branch error in the associated quadrature. Counting q- and p-boundary crossings therefore gives a software-level indicator of how often the fast path leaves the residual on the wrong side of the correctable cell. It does not capture the full logical channel, correlations across repeated correction rounds, leakage, auxiliary state noise or outer-code behavior. Those omissions are why the metric is reported as `final_ler_mean` rather than as logical-channel fidelity or hardware logical-error rate.

The relation between the reported proxy and the channel-style surrogate can be made explicit. For round t , define

$$X_t = \mathbb{1}\{|r_{q,t}| > \lambda/2\}, \quad Z_t = \mathbb{1}\{|r_{p,t}| > \lambda/2\},$$

where $r_{q,t}$ and $r_{p,t}$ are the components of r_t^+ in Eq. (7). The reported run-level proxy is

$$\text{final_ler} = \frac{1}{T} \sum_{t=1}^T (X_t + Z_t).$$

The union event used in the surrogate channel table is

$$U_t = \mathbb{1}\{X_t \vee Z_t\} = X_t + Z_t - X_t Z_t.$$

Thus $\mathbb{E}[U_t]$ and $\mathbb{E}[X_t + Z_t]$ are linked but not identical: a simultaneous q/p crossing contributes once to U_t and twice to `final_ler_mean`. In expectation, $p_{\text{any}} \leq \mathbb{E}[X_t + Z_t] \leq 2p_{\text{any}}$. This is why the manuscript reports `final_ler_mean` as the primary software metric and uses the surrogate channel language only as an interpretive bridge from boundary events to Pauli-like categories. A genuine finite-energy logical-channel fidelity would require propagating the corrected approximate-GKP state or an equivalent logical process map, not only counting residual boundary events [31, 32, 33].

3.4 Boundary sensitivity and surrogate channel language

The half-lattice boundary also gives a compact theoretical bridge between residual-scale metrics and channel-style language. If a residual quadrature is approximated locally as $r \sim \mathcal{N}(0, \sigma_r^2)$, the probability of a single-quadrature boundary crossing is

$$p_{\text{cross}}(\sigma_r) = \text{erfc}\left(\frac{\lambda/2}{\sqrt{2}\sigma_r}\right). \quad (9)$$

Assuming independent q and p residuals with the same local scale gives

$$p_{\text{any}} = 1 - (1 - p_{\text{cross}})^2. \quad (10)$$

This calculation is not a finite-energy logical-channel simulation. It is a diagnostic bridge: it explains why residual variance near the decision boundary can change a boundary-crossing metric sharply, and why `final_ler_mean` and residual MSE are related but not interchangeable.

3.5 Metric interpretation

The notation above is tied to a concrete software implementation, but only at the level needed for a scoped protocol claim. The codebase uses the same GKP lattice constant $\lambda = \sqrt{2\pi}$ for syndrome wrapping, half-lattice logical boundaries and fast-loop histogram limits. The measurement model includes finite-energy and readout terms through a compact configuration, and the fast-loop emulator records residual accumulation with the same half-lattice logical-error rule. These anchors support the statement that `final_ler_mean` is a protocol-defined logical-error proxy for the current controlled simulation setting. They do not turn the benchmark into a complete finite-energy GKP physical model, a wrapped-posterior decoder comparison, a holdout drift distribution or a hardware measurement. The current supporting-data checks are consistency checks between manuscript tables, code paths and generated figures; they are not physical validation.

4 Method

4.1 Overview

The method decomposes drift-adaptive correction into three modules (Fig. 1). The fast path applies the active affine correction. The teacher or calibration module estimates a low-dimensional noise state from recent histograms. The runtime controller stages, clips, smooths and commits candidate parameters into the bank consumed by the fast path.

Each module has a distinct scientific and engineering role. The fast path is the object that must eventually meet a hard timing budget, so it is deliberately limited to clipping, fixed-point arithmetic, a 2×2 matrix-vector operation and status logging. The slow estimator is the object that absorbs nonstationarity, so it is allowed to use windowed histograms, teacher features or statistical summaries unavailable to a single-shot combinational datapath. The runtime controller is the safety interface between them: it prevents an unstable estimate from becoming an active correction until the candidate surface has passed clipping, quantization and commit checks. This separation lets the paper compare estimator families without changing the timing-critical operation being validated.

4.2 Design rationale and testable predictions

The design is useful only if it makes stronger claims easier to test. The central hypothesis is not that an affine rule is a complete GKP decoder, but that the adaptive state of a drift-aware physical

correction layer can be compressed into a small surface whose online arithmetic is deterministic. This compression creates four testable predictions. First, if drift mainly changes local bias, scale or orientation, recent syndrome statistics should improve the committed (K, b) surface and lower residual-boundary events relative to a static or slowly mismatched affine reference. Second, if the estimator is kept outside the per-shot path, the online cost should scale with the affine datapath rather than with posterior branch count or neural-network depth. Third, if the same committed state is logged in software and hardware, the future board implementation should be checkable by source-vs-board vectors rather than by only aggregate logical-error rates. Fourth, if the local single-branch assumption fails, the method should fail in identifiable ways: posterior mass should split across lattice branches, half-lattice crossings should increase, or stale parameters should dominate the commit-lag sweep.

These predictions determine how the experiments are organized. The main controlled simulation benchmark tests the first prediction under four controlled drift families. The analytical cost model and Q4.20 parity check test the second prediction at the software and arithmetic-contract level. The runtime counters and validation-contract figure make the third prediction observable, although board measurements are still absent. The oracle-affine, wrapped-Gaussian, holdout-stress and commit-lag diagnostics address the fourth prediction by showing when local affine tracking is plausible and where it should be distrusted. This rationale is the reason the manuscript reports multiple bounded metrics instead of converting them into a single headline score.

4.3 Algorithmic protocol

The protocol-level object is a sequence of committed affine surfaces rather than a single neural prediction. At each slow-loop update, the implementation uses the recent syndrome window $H_{t-c+1:t}$, its time differences and any available teacher or statistical state to propose a candidate $(K_t^{\text{cand}}, b_t^{\text{cand}})$. The candidate is clipped, optionally smoothed, checked against fixed-point and saturation limits, and staged in an inactive parameter bank. A commit event makes the candidate the active surface only at a safe epoch boundary. Between two such commits, every shot uses the same active surface through Eq. (1).

This organization makes the estimator interchangeable but keeps the fast-loop contract fixed. A teacher-anchored residual branch, a statistical calibration rule or a future estimator may differ in how it proposes $(K_t^{\text{cand}}, b_t^{\text{cand}})$, but the logged runtime state is the same: active K, b , staged K, b , clipping and saturation status, commit acknowledgement, stale-parameter counters and the resulting residual boundary events. These fields are the minimum information needed to reproduce the software protocol and to compare a future board implementation against the same source trajectory. The present manuscript evaluates this protocol in software; it does not report board commit latency or source-vs-board agreement.

4.4 Affine fast path

Given the measured syndrome s_t , the fast path computes

$$s_t^{\text{clip}} = \text{clip}(s_t, -s_{\text{max}}, s_{\text{max}}), \quad (11)$$

$$\Delta_t^{\text{raw}} = K_t s_t^{\text{clip}} + b_t, \quad (12)$$

$$\Delta_t = Q(\text{clip}(\Delta_t^{\text{raw}}, -\Delta_{\text{max}}, \Delta_{\text{max}})), \quad (13)$$

where $Q(\cdot)$ denotes fixed-point quantization. A hardware implementation can represent K_t , b_t , clipped syndromes and residual corrections in a selected $Q_{m,n}$ format, expose saturation flags, and

route the active parameter bank through a small matrix-vector datapath. The fast-path latency model is correspondingly simple:

$$L_{\text{fast}} = L_{\text{wrap}} + L_{\text{matvec}} + L_{\text{clip}} + L_{\text{bank}}. \quad (14)$$

This module is needed because the per-shot QEC path must remain small enough for deterministic execution and independent validation. Its advantage is not that affine decoding is globally optimal, but that all more complex estimation can be moved outside the fast timing boundary.

The latency expression is meant as a design decomposition rather than a board measurement. A fully parallel implementation can assign four fixed-point multipliers to the 2×2 product and reduce the matrix-vector stage to a short adder tree plus output registers. A resource-minimal implementation can reuse fewer multipliers across several cycles, increasing L_{matvec} while reducing DSP use. Both realizations execute the same committed map and can be checked by the same tuple stream. This is the main hardware-facing benefit of the formulation: latency and area can be traded by changing the arithmetic schedule without changing the estimator, the scientific metric or the source trace.

4.5 FPGA-facing datapath contract

The hardware-facing object in this study is a datapath contract, not a measured device implementation. The contract has five parts: a wrapped two-quadrature syndrome input; an active bank containing K_t , b_t and clip thresholds; a fixed-point affine matrix-vector operation; a corrected displacement output; and status flags for clipping, overflow, stale parameters and bank commits. This decomposition is useful because every fast-path output can be checked against a software reference with explicit input vectors and parameter-bank state. It also makes the negative comparison precise: a per-shot wrapped-posterior decoder must evaluate multiple branch hypotheses inside the timing boundary, whereas the proposed path evaluates one clipped affine rule and leaves posterior or neural estimation to the slow loop. In register-level terms, the fast path is intentionally small. A minimal implementation needs four fixed-point registers for K_t , two for b_t , clip-threshold registers, an active-bank selector, a commit acknowledgement flag and overflow or saturation status flags. The online datapath then performs a bounded sequence of operations: read the active bank, clip the wrapped syndrome, compute the two-row affine product, add the bias, quantize the displacement and emit status bits. No histogram update, CNN inference, posterior branch search or parameter fitting is required on the correction cycle. This is the practical meaning of the arithmetic-footprint comparison: the adaptive computation may be expensive, but it is moved behind a staged commit interface rather than placed in the per-shot correction path.

The same decomposition defines how the design would be tested on an FPGA. A source-level test vector is the tuple

$$(s_t^{\text{clip}}, K_t, b_t, \Delta_t, \text{flags}_t, \text{bank}_t),$$

recorded at the shot boundary before and after a commit. A board run would be asked to replay this tuple stream with the same fixed-point format and to match the source correction and status bits within the declared quantization tolerance. This requirement is narrower than an end-to-end decoder benchmark but stronger than a schematic hardware claim: it specifies exactly which registers, bank selectors, status flags and arithmetic outputs must agree before the software contract can support a device-level claim.

This source-vs-board criterion is also the reason the manuscript keeps the fast-path object affine. A neural or wrapped-posterior per-shot datapath would require validating many hidden activations, branch scores or nonlinear approximations in addition to the final correction. The affine contract

reduces the future device check to six active parameters, two clipped inputs, two correction outputs and a small set of flags. The reduction does not make the present study a hardware experiment, but it makes the missing hardware experiment precisely specified.

The contract is deliberately phrased in the same measurement language used by real-time decoder studies. A hardware claim would require measured latency, resource and power numbers, bit-accurate fixed-point agreement and a logged source-vs-board trace. Before those measurements exist, the available support is software-side and analytical: the affine operation count is 4 multiplications and 4 additions with 6 stored state scalars; Q4.20 software emulation shows no residual-boundary crossing-rate change in the controlled sample set; and software runtime counters expose commits, fast-cycle violations, overflow and saturation. These checks make the fast path a concrete candidate for a future FPGA experiment, but they do not substitute for synthesis or board data.

4.6 Teacher-anchored residual calibration

The teacher estimator maps a recent histogram window $H_{t-c+1:t}$ to an effective state $\hat{\theta}_t^{\text{teacher}}$, then to teacher parameters through the affine mapper:

$$(K_t^{\text{teacher}}, b_t^{\text{teacher}}) = \text{ParamMapper}(\hat{\theta}_t^{\text{teacher}}). \quad (15)$$

The learned branch predicts only a bounded residual bias term,

$$\hat{\delta b}_t = f_\phi(H_{t-c+1:t}, \Delta H_{t-c+2:t}, z_t^{\text{teacher}}), \quad (16)$$

followed by clipping and smoothing:

$$b_t = \text{EMA}\left(b_t^{\text{teacher}} + \text{clip}(s_b \hat{\delta b}_t, -b_{\max}, b_{\max})\right), \quad K_t = K_t^{\text{teacher}}. \quad (17)$$

The motivation is to let the teacher define a stable physical calibration surface while the residual branch handles systematic mismatch left by the teacher. The design is intentionally conservative: the neural network never directly replaces the fast GKP correction law.

The residual branch used in the reported experiments is deliberately small and runtime-aligned (Table 10). Its inputs are the same window-level quantities available to the slow loop, and its labels are residual affine-bias corrections rather than direct logical decisions. This choice reduces leakage from the evaluation objective into the per-shot path: the model learns how to correct the teacher’s committed bias surface, while the fast loop still receives only (K, b) .

4.7 Stage-and-commit runtime contract

The slow loop does not mutate active parameters immediately. It stages a candidate (K, b) in an inactive bank, validates coefficient bounds, quantization error and saturation risk, and commits at a safe epoch boundary. The contract exposes observable quantities: update cadence, stale-parameter counters, commit acknowledgements, rollback events, overflow and saturation. In this study, these quantities are evaluated only in controlled simulation and supporting software-runtime checks. Board-level commit latency, source-vs-board agreement and resource use remain unmeasured.

4.8 Supplementary statistical calibration

The same affine contract can host non-neural estimators. A scoped statistical-calibration branch computes a bias correction directly from recent syndrome statistics and emits parameters through the same clipping and commit path. This branch is useful scientifically because it tests whether the

affine calibration contract, rather than the CNN architecture alone, explains the observed gains. It is reported as a supplementary analysis in this manuscript and remains separate from the main controlled simulation comparison.

5 Experimental Design

5.1 Claim-to-test map

We evaluate four questions. First, can a slow-loop residual branch lower the protocol-defined `final_ler_mean` of a deterministic affine GKP fast path under controlled drift? The predeclared five-mode controlled simulation benchmark addresses this question. Second, is the effect better understood as adaptive calibration of a low-dimensional affine surface than as an irreplaceable neural architecture? Feature ablations, six-seed mechanism summaries and the supplementary statistical-calibration branch address this question. Third, when is the affine rule a meaningful local approximation, and when should a wrapped posterior be distrusted? Oracle-affine and wrapped-Gaussian checks address this question. Fourth, is the committed affine path small enough to be a concrete fixed-point and source-vs-board target? Operation counts, software fixed-point parity and the hardware-measurement plan address this question, while board measurements remain outside the present study.

The evaluation keeps these questions separate. A successful result on one metric is not treated as support for several stronger endpoints. The main benchmark can support a controlled drift ranking, but not logical-channel fidelity, statistical significance, real-time FPGA latency or end-to-end surface-GKP overhead. Those stronger claims require the separate evidence classes listed in Tables 3 and 32.

5.2 Predeclared simulation benchmark

The main experiment uses a predeclared four-scenario, five-mode controlled simulation protocol. The five main modes are an extended Kalman filter, an unscented Kalman filter, a constant-mean affine baseline, an RLS residual- b baseline and the teacher-anchored hybrid residual- b branch. All reported main rows use the same predeclared scenario matrix and paired-seed/repeat setting. This benchmark tests adaptive calibration laws under matched synthetic drift trajectories. It does not test a real FPGA board, a default embedded runtime, or unseen drift families. It also does not compare against a tuned finite-energy nearest-lattice or sequence-level wrapped-posterior GKP decoder; the nearest-syndrome, oracle-affine and wrapped-Gaussian rows are diagnostic checks for the affine runtime-calibration protocol.

5.3 Comparator ladder and stress-test roles

The comparison design is intentionally layered because no single baseline can evaluate all scientific questions posed by a drift-adaptive GKP fast path. The main controlled simulation table compares deployable calibration laws under the same runtime interface. The controlled oracle and wrapped-Gaussian checks ask whether the affine form is locally meaningful when the effective state is known. The holdout stress tests ask whether update latency, discontinuities or faster-than-window drift can erase the benefit of a better affine surface. The statistical-calibration extension asks whether the contract is tied to a particular neural estimator. Table 11 summarizes these roles before the results are reported.

5.4 Ablation and mechanism materials

The ablation materials remove or alter histogram-delta inputs and teacher-side features. The six-seed mechanism pack summarizes descriptive cross-seed behavior and a lower-clip intervention. These materials are included to explain the main result conservatively. They are not designed to establish a causal mechanism.

5.5 Hardware-validation protocol

Hardware-facing validation is specified as a future-measurement protocol rather than as an available result. A board-level validation study would report the board platform, bitstream or firmware hash, AXI/DMA interface, fixed-point format, fast-path latency distribution, slow-loop update time, commit acknowledgement latency, resource use, power and source-vs-board numerical agreement. Until those measurements exist, hardware is treated as a validation target rather than as a completed result.

6 Results

The Results first report the predeclared controlled simulation ranking used for the main performance statement. They then add paired uncertainty summaries and two scenario-level paired-repeat comparisons, followed by controlled diagnostics that separate estimator limits, drift sensitivity and implementation feasibility. The final result layers report analytical cost, fixed-point software parity and runtime-discipline counters. Read together, these data support an affine-calibration systems method under controlled simulation while leaving finite-energy channel fidelity, expanded statistical robustness and hardware execution as open validation requirements.

6.1 Hybrid residual calibration ranks lowest in the predeclared simulation set

Under the predeclared four-scenario protocol, the hybrid residual branch has the lowest mean `final_ler_mean` in all four tested scenarios, while UKF is the runner-up in each case (Table 12). Relative to UKF, the hybrid branch reduces mean `final_ler_mean` by 1.75% in static-bias drift, 2.89% in linear ramp drift, 2.80% in the step-change scenario and 1.85% in periodic drift, for an average reduction of 2.32% across the tested set. This is the direct UKF-versus-hybrid comparison under the shared protocol; the later diagnostic baselines test oracle access, posterior branches, holdout stress and estimator interchangeability. The paired raw rows are directionally consistent across the two available repeats in each scenario (Tables 13, 14 and 15). The result is consistent with a protocol-level improvement from slow-loop residual calibration in the tested scenarios. The current run contains two repeats per scenario and mode. The corresponding standard deviations, paired deltas and paired-bootstrap envelope are therefore reported only as descriptive uncertainty markers; they should not be interpreted as confidence intervals, standard errors or hypothesis-test evidence.

Two scenarios also have completed formal-length paired-interval checks with 12 paired repeats. For `static_bias_theta` and `linear_ramp`, the UKF-minus-hybrid `final_ler_mean` intervals remain positive under both paired- t and bootstrap summaries (Table 44). These rows strengthen the scenario-level evidence for the two completed drift families, but they do not yet support pooled, all-scenario or holdout claims.

To make the ranking interpretable as a magnitude statement rather than only a winner label, we also report the absolute UKF-minus-Hybrid margin and its scale relative to the larger of the two




Figure 2: Main controlled simulation comparison. The plot visualizes the existing `final_ler_mean` means from Table 12; lower is better. Error bars show descriptive standard deviations across the two available repeats per scenario and mode. The plot supports a predeclared four-scenario controlled simulation ranking only. Paired repeat deltas are provided in Table 13, the descriptive resampling envelope is provided in Table 14, and data-backed advantage margins are provided in Table 15; all are bounded by the accompanying supporting data. Inferential confidence intervals, formal hypothesis tests and unseen drift families remain open validation requirements.

reported descriptive standard deviations (Table 15). This ratio is a supporting-data readout over the current paired rows, not a significance statistic. It reports that the source-row margin exceeds the larger descriptive SD in these rows, without implying statistical separation and while preserving the same $n = 2$ statistical boundary as the main table.

A repeat-expanded evaluation protocol is specified for future stronger statistical claims, but it is not counted as performance evidence in the present Results. The protocol keeps the five-mode table as the anchor, requires additional UKF-versus-Hybrid paired rows before all-scenario inferential wording, and reserves random-walk, burst/reset and faster-than-window families for holdout validation. The present manuscript still does not report formal hypothesis tests or an executed holdout benchmark.

We also ran a shortened UKF-versus-Hybrid feasibility check over all four predeclared scenarios with paired seeds and two repeats. The check completed all requested rows and preserved the same direction as the main benchmark, but the shortened timing configuration makes it a field-availability check rather than additional performance evidence (Table 16).

6.2 Controlled nearest-syndrome, oracle and wrapped-Gaussian checks separate estimator and decoder limits

These controlled baselines answer a narrow technical question: whether the affine form is obviously dominated by a hard nearest-syndrome correction or a known-state wrapped posterior rule in a local one-step setting. They are sanity checks for the fast-path approximation, not substitutes for a formal known-noise or nearest-lattice benchmark.

To isolate whether the affine fast path is limited by the affine form itself, by the upstream estimator or by the absence of a posterior branch model, we added a small controlled local-Gaussian analysis. The analysis samples 120,000 one-step errors for four representative drift states, compares a hard nearest-syndrome correction, a fixed nominal affine rule and an oracle affine rule that is given the state $(\sigma, \mu_q, \mu_p, \vartheta)$, and adds one-step wrapped-Gaussian posterior-mean and MAP branch baselines.

The result is deliberately narrower than the main benchmark. The oracle affine rule reduces residual MSE relative to the fixed nominal affine rule in all four controlled states, with the largest reduction in the post-step state (Table 17). The nearest-syndrome row is much weaker in residual MSE because it applies the noisy wrapped syndrome directly as a hard correction, making measurement noise and branch ambiguity visible. The wrapped-Gaussian posterior mean is slightly better than oracle affine only in the static state; in the ramp, step and periodic states it has higher MSE, and the MAP branch rule is weaker still. This mixed outcome is important. It shows that known-state affine correction captures a real local modelling advantage, while a nearest-syndrome or wrapped-posterior comparator cannot be treated as solved by simply adding a hard branch or branch prior to the one-step setting. A fully tuned sequence-level wrapped or nearest-lattice baseline, a holdout drift set and `final_ler_mean`-level scoring remain required before making stronger decoder-comparison claims.

6.3 A residual-boundary channel surrogate clarifies what `final_ler_mean` does and does not measure

The half-lattice crossing rule can be written in channel language without claiming a full logical-channel reconstruction. For each controlled local-Gaussian state, we decomposed residual boundary events into q-only, p-only, simultaneous q/p and no-crossing outcomes. This gives a Pauli-channel-style surrogate

$$p_I^{\text{surr}} = 1 - p_{\text{any}}, \quad p_q^{\text{surr}}, \quad p_p^{\text{surr}}, \quad p_{qp}^{\text{surr}},$$

where p_{any} is the probability that the residual crosses at least one half-lattice boundary. If this decomposition is treated only as a qubit Pauli-channel surrogate, it induces $F_{\text{avg}}^{\text{surr}} = (1 + 2p_I^{\text{surr}})/3$. The manuscript does not use this quantity as a finite-energy GKP logical-channel fidelity; it is included in the supporting data only to make the relationship between residual-boundary scoring and fidelity language explicit. The point of the table is therefore preventive as much as comparative: it shows exactly where the protocol-defined LER proxy stops and where a genuine finite-energy channel analysis would have to begin.

Table 18 reports the safer quantity, p_{any} . This union rate is related to, but not identical with, the manuscript `final_ler_mean` definition: `final_ler_mean` counts q- and p-boundary crossings separately, so a simultaneous q/p event contributes twice, whereas p_{any} asks whether at least one Pauli-like residual-boundary event occurred in the sample. The static and ramp states have no observed affine crossings in 120,000 controlled samples, whereas the post-step and periodic states expose low but nonzero residual-boundary events. The wrapped-Gaussian MAP branch rule has higher surrogate crossing probability in all four states, again showing that branch-aware posterior

logic needs a tuned sequence-level protocol before it can serve as a stronger comparator. The table supports only a residual-boundary event decomposition; it does not estimate leakage, finite-energy envelopes, auxiliary-state noise, correlated rounds or process fidelity.

The same supporting data can also be read at the method level (Table 20). This aggregation is useful because it separates the question “which method has the larger residual-boundary channel surrogate?” from the per-state event decomposition. Across the four controlled local-Gaussian states, the fixed and oracle affine rows have the lowest mean p_{any} , while the wrapped posterior MAP row has the largest worst-state surrogate crossing rate. This is still a lattice-level residual-boundary sanity summary, not a finite-energy channel measurement.

To move one step closer to finite-energy channel language without claiming tomography, we also ran a simplified finite-squeezing measurement-channel sanity check. This check samples the same controlled drift states as the residual-boundary analysis, adds a scalar finite-squeezing term Δ and an effective measurement-efficiency contribution to the wrapped syndrome, and then compares three correction rules: hard nearest-syndrome correction, the fixed affine fast path and an oracle affine rule. The largest toy-channel stress appears at $\Delta = 0.34$, where hard nearest-syndrome correction has a mean p_{any} of 1.002×10^{-3} , while the fixed affine row remains at 1.625×10^{-4} in the same toy model (Table 21). This is useful as a model-consistency check because it exercises the same half-lattice decision surface after adding an explicit finite-squeezing noise scale. It is still not a calibrated finite-energy GKP logical-channel simulation, process fidelity estimate or hardware result.

6.4 Sequence-level controlled baselines expose posterior-branch risk under drift

We next repeated the nearest-syndrome, oracle and wrapped-Gaussian comparison over short drift sequences rather than isolated samples. Each scenario used 384 independent sequences of 512 steps, the same local-Gaussian state families as the main controlled analysis, and the same half-lattice residual-boundary proxy used in the manuscript metric. The sequence result preserves the one-step conclusion that known-state oracle affine lowers residual MSE relative to a fixed nominal affine rule, while the direct nearest-syndrome reference gives higher crossing proxies in all four scenarios. It also exposes a practical risk of posterior-branch rules: the wrapped-Gaussian posterior mean and MAP references produced higher sequence-level crossing proxies in the ramp, step and periodic states (Table 22).

This result does not establish that affine decoding is globally superior to wrapped decoding. It shows that a branch-aware posterior comparator must be implemented and tuned as a sequence-level decoder, not inserted as a naive one-step branch rule. The same caution applies to nearest-lattice-style references: the present hard nearest-syndrome row is a stress reference, not a formal nearest-lattice decoder. For this manuscript, the result strengthens the engineering rationale for a calibrated affine fast path while keeping the stronger-baseline requirement open.

6.5 Holdout drift stress tests expose adaptation space and lag sensitivity

We next tested three drift patterns that were not part of the main four-scenario comparison: a clipped random walk in the effective noise state, two burst/reset intervals, and an oscillatory drift faster than the slow-loop commit window. The analysis used the same local-Gaussian displacement model and the same 384-by-512 sequence shape as the controlled sequence study. It compared a fixed nominal affine rule, an oracle affine rule, a lagged-affine reference with a 64-step commit interval and a 32-step lag, and the same one-step wrapped-Gaussian posterior references.

The holdout stress test is informative because the half-lattice crossing proxy is rare in these controlled samples. The affine rows have the same observed crossing rate within each holdout

scenario, so residual MSE is the more sensitive stress-test statistic. The same supporting data also report a residual-boundary Pauli-style identity surrogate; the oracle-affine values remain near unity because crossings are rare, but this is still not a finite-energy logical-channel fidelity. Oracle affine gives the lowest residual MSE in all three holdout families (Table 23), showing that a correctly updated affine surface still has adaptation headroom outside the main drift set. The lagged-affine row is deliberately mixed: it approaches oracle behavior under random-walk drift, but it is worse than the fixed rule under burst/reset and faster-than-window oscillation. This failure mode is a useful constraint on the method. It indicates that the affine fast path is not the limiting idea by itself; stale parameter commits can erase the adaptation benefit under rapid or discontinuous drift.

The result does not prove generalization of the trained residual branch, because the lagged row is a known-state slow-commit reference rather than the CNN calibration module. It also does not provide confidence intervals or a formal expanded benchmark. Its role is narrower: it adds an unseen-drift stress layer that makes the proposed fast-path contract more falsifiable and turns drift-tracking latency into a measurable design variable.

The same diagnostics can be read as a local-validity check for the affine approximation (Table 24). The short-sequence rows ask whether an oracle affine surface has residual-MSE headroom over a fixed nominal surface, and whether a naive wrapped-Gaussian posterior mean reduces the half-lattice crossing proxy. The holdout rows ask the same MSE headroom question under unseen drift families and add a stale-commit risk term, measured as lagged-affine residual MSE minus oracle-affine residual MSE. This view is intentionally diagnostic: it supports the claim that affine calibration has a local validity regime and identifiable failure modes, while leaving a formal nearest-lattice, known-noise or wrapped-decoder benchmark as a future requirement.

6.6 Commit-lag sweeps turn stale parameters into a design variable

The fixed 32-step lag in Table 23 is only one point on a slow-loop design surface. We therefore swept the commit lag over 0, 8, 16, 32, 64 and 128 simulation steps while keeping the commit interval fixed at 64 steps. The sweep reuses the three holdout drift families and shared sampled errors within each family, so it is a controlled stale-parameter diagnostic rather than a new benchmark.

The resulting pattern is deliberately not over-simplified (Table 25). Under random-walk drift, longer lag monotonically increases residual MSE while still remaining below the fixed nominal affine reference. Under burst/reset and faster-than-window drift, the lag surface is non-monotone because the stale state can occasionally land closer to a later reset or oscillation phase than a recently committed state. This is precisely why commit latency should be measured rather than inferred: the fast affine surface remains cheap, but the slow-loop commit policy is part of the method’s empirical contract.

6.7 Analytical operation counts quantify the fast-path cost argument

The cost advantage of the affine fast path is an analytical datapath argument, not yet a measured FPGA-timing result. The count assumes a two-quadrature syndrome input, a 2×2 affine matrix, one bias vector, a single active parameter bank and one correction per fast-loop shot. It counts arithmetic and small-state storage on the latency-critical path, while excluding the slow-loop estimator, memory-controller implementation details and clock-frequency closure. Under this model, the affine rule uses one branch, four multiplications, four additions and no exponential or division operations (Table 26). A 3-by-3 wrapped-Gaussian MAP reference increases the multiplication count by 12.25-fold and the stored state by 5.5-fold, while a posterior-mean reference requires softmax-like nonlinear operations in addition to a larger arithmetic footprint. This comparison supports the

Software validation contract: low arithmetic, fixed-point parity and observable counters




Figure 3: Software validation contract for the affine fast path. Left: analytical multiplication and addition counts compare the affine rule with one-step wrapped-Gaussian references. Middle: Q4.20 software fixed-point emulation shows small p99 correction differences and no observed residual-boundary crossing-rate change in the controlled sample set. Right: software-in-the-loop runtime counters make overflow and fast-cycle violations visible under the current protocol. The figure summarizes data-backed software checks; it is not FPGA synthesis, timing closure, power/resource measurement, source-vs-board agreement or board commit latency evidence.

design choice to keep posterior or learned estimation outside the latency-critical correction path. The cost model is a candidate-selection argument for the future hardware path, not a hardware measurement. It does not claim timing closure, resource use, power or source-vs-board agreement. Together with the fixed-point and runtime-counter checks below, it provides the software evidence currently available for the datapath contract in Table 9.

6.8 Software fixed-point emulation bounds numerical degradation

The affine fast path is intended to be representable as fixed-point arithmetic, so the manuscript also checks numerical degradation before any board claim is made. We compared the floating-point affine runtime with the same runtime under Q4.20 fixed-point emulation on 80,000 controlled local-Gaussian samples per state. The largest correction difference was 1.6×10^{-6} , no scenario changed the residual-boundary crossing rate, and no fixed-point quantization saturation occurred (Table 27). This supports the fixed-point feasibility of the affine datapath under the tested parameter range. It does not establish board timing, source-vs-board agreement, resource use or power.

6.9 Runtime counters make the software contract observable

The same controlled simulation comparison records runtime counters for commit activity, cycle violations, overflow and saturation. These counters do not measure board latency or hardware reliability, but they make the stage-and-commit contract observable in the software protocol. Across the scenario-level comparison rows, each mode applied about 900 commits on average, slow-update violations were zero, correction saturation was zero, and the observed overflow rate was approximately 2.5×10^{-3} , dominated by histogram-input saturation (Table 28). This result supports the narrower claim that the runtime contract exposes measurable software counters; it does not establish board commit timing, rollback reliability or source-vs-board agreement.




Figure 4: Ablation and descriptive mechanism evidence. Left: average LER change relative to UKF for the scoped feature ablation table. Right: lower-clip intervention effect across the six-seed mechanism pack. The figure supports feature sensitivity and a cautionary intervention reading; it does not establish a causal mechanism, teacher necessity or a successful mitigation.

6.10 Feature ablations constrain the mechanism story

The feature and teacher ablation table shows that input choices materially change performance in the tested scenario set (Table 29). Removing histogram deltas degrades performance relative to the full hybrid model, which supports the use of recent histogram dynamics. However, the no-teacher-parameters row has the best four-scenario average. The result therefore does not support a simple claim that teacher parameters are necessary or always beneficial. It supports a weaker but more credible interpretation: the residual calibration design is sensitive to feature composition, and the current mechanism remains descriptive.

6.11 Cross-seed intervention evidence remains descriptive

The six-seed mechanism pack indicates that an instability pattern appears across multiple seeds, but the tested lower-clip intervention is mixed and mostly harmful (Table 30). This result is valuable because it prevents an over-simple mechanism story. Residual amplitude and committed bias dynamics may participate in failures, but the current intervention does not show that lowering the residual clip is a general fix.

6.12 Supplementary statistical calibration supports the contract

The statistical-calibration rows are reported under a separate supplementary protocol rather than under the predeclared main ranking protocol (Table 31). They use the same affine runtime interface but answer a different question: whether a non-neural estimator can populate the committed (K, b) surface at all. They should therefore be read as contract evidence, not as evidence that the neural




Figure 5: Supplementary statistical-calibration analysis. The plot compares UKF, hybrid-*b* and the statistical-calibration values already listed in Table 31. The result supports the broader affine calibration contract, but it must remain separate from the main comparison because it was not generated under the same predeclared ranking protocol; the absolute values should not be interpreted as a matched win over the main hybrid branch.

branch was beaten in a fair main-table comparison. The absolute values in this extension are not comparable to the main ranking because candidate generation and protocol generation differ. A fully matched comparator claim would require a single predeclared protocol that fixes candidate generation, selection, repeats, uncertainty handling and reporting before seeing the outcomes.

6.13 Required hardware and deployment measurements

The present manuscript does not contain real-board execution results. The available deployment-adjacent material supports only one narrow runtime statement: selected preserved `.tflite` artifacts executed in an isolated software runtime. Board execution, timing, resource, power and source-vs-board measurements were outside the scope of this study. Table 32 defines the hardware measurements required for a future implementation-stage evaluation.

7 Discussion

The present results support a specific systems claim for drift-adaptive GKP decoding: nonstationary syndrome statistics can be converted into updates of a physical-layer affine correction surface while the per-shot correction remains deterministic. The main evidence is the predeclared ranking of the hybrid residual branch under a controlled simulation protocol. Scoped ablations and the supplementary statistical-calibration analysis sharpen the interpretation. Teacher features, histogram dynamics and estimator choice affect the result, but the object that persists across these

variants is the committed affine surface consumed by the fast path. This places the work between soft-information GKP decoding, which exploits analog syndrome structure, and real-time hardware decoding, which requires bounded latency and verifiable datapaths. The distinctive design choice is the separation of roles: the slow loop absorbs nonstationary statistics, while the correction cycle evaluates only a small affine matrix-vector update.

The supplementary statistical-calibration analysis reinforces this reading from the opposite direction. If a simple non-neural estimator can be strong under the same runtime contract, the scientific object is the contract itself: recent syndrome statistics update a low-dimensional affine surface that the fast path can consume safely. The CNN branch is one useful slow-loop estimator, not the only possible estimator and not yet a final comparator.

Several alternative explanations are therefore constrained, but not eliminated. The paired rows and advantage-margin readout show that the UKF-versus-hybrid direction is consistent in the available repeats, but the two-repeat design is still too small for an all-scenario inferential statistical claim. The completed 12-paired-repeat `static_bias_theta` and `linear_ramp` checks are important scenario-level exceptions: both give positive paired- t and bootstrap interval lower bounds for UKF-minus-hybrid `final_ler_mean` deltas. They do not, by themselves, justify pooled or all-scenario significance language. The ablation and statistical-calibration results make a pure “CNN wins because it is neural” interpretation too strong: feature design, histogram dynamics and the affine commit contract remain part of the effect. The residual-boundary channel surrogate also rules out a purely notational fidelity claim, because it states exactly which q/p half-lattice crossing events are being counted and which finite-energy channel quantities remain outside the analysis. Finally, the oracle and wrapped-Gaussian checks indicate that a local posterior-style rule is not an automatic replacement for the affine contract under drift, while also showing that stronger known-noise and sequence-level baselines remain necessary before making a broad decoder-comparison claim.

The supported claim is therefore an interface-level design trade-off, grounded in three linked but non-substitutable tests. First, the software benchmark tests whether recent syndrome statistics can improve the protocol-defined `final_ler_mean` of a deterministic correction surface under drift. Second, the residual-boundary channel surrogate tests whether that same improvement can be translated into q/p event language without upgrading it to finite-energy logical-channel fidelity. Third, the implementation-feasibility checks test whether the online rule remains small enough to be a plausible real-time datapath: latency is addressed structurally by placing estimation in the slow loop and leaving the fast path as $\Delta_t = K_t s_t + b_t$, while cost and verifiability are supported by the analytical operation-count model, fixed-point software parity and the small source-vs-board target implied by a clipped affine datapath. These points motivate hardware validation, but they are not hardware results: board timing, resource use, power, commit reliability and source-vs-board vectors remain to be measured.

The numerical scale of adjacent work clarifies the comparison rather than changing the supported claim. Calibration-conditioned and learned decoders report real-processor LER reductions, large calibration data sets and microsecond-scale neural latency anchors; real-time hardware decoders report nanosecond-to-sub-microsecond decoding or feedback paths, lookup-memory compression, device scale, area, power and cycle budgets. Those numbers are not baselines for the present software study. They set two requirements for a future stronger version of this work. First, the affine fast path must be measured on a device against the same timing, resource, power and source-vs-board standards used in hardware-decoder papers. Second, the scientific value of keeping learning in the slow loop must be judged by whether it preserves drift adaptation while avoiding a per-shot neural datapath. The present manuscript supports that architecture-level hypothesis, not the hardware claim itself.

This distinction is the main comparative trade-off of the present formulation. Relative to analog

GKP and surface-GKP decoders, the method does not attempt to replace the outer-code likelihood engine. It supplies a physical-layer calibration surface whose state can be logged, clipped and committed. Relative to calibration-conditioned neural decoders, it does not place the learned model on the per-shot path; the online computation remains a deterministic affine correction even when the committed parameters were produced by a CNN, a teacher anchor or a statistical estimator. Relative to hardware-tailored FPGA decoders, it does not yet provide measured device latency, but it does provide a smaller and more directly replayable target for future source-vs-board validation. The present data therefore support an interface-level design trade-off rather than a single universal advantage: controlled-drift adaptation is paired with a bounded arithmetic footprint, fixed-point-checkable numerical behaviour and a traceable parameter-bank contract, while measured hardware advantage remains untested. The same affine contract also separates the future latency-area trade-off from the estimator design: parallel multipliers and serialized multipliers implement the same committed map and can be validated against the same source trace. The verification target is smaller for the same reason the online datapath is smaller. A board implementation would need to reproduce six affine parameters, two clipped syndrome inputs, two correction outputs and a bounded set of status flags at each shot boundary. A per-shot neural or wrapped-posterior datapath would also require validating hidden activations, branch scores or nonlinear approximations. The expected advantage is therefore not only fewer arithmetic operations, but a narrower object for bit-accurate replay and failure analysis. This reading also explains why the manuscript keeps four metrics side by side instead of collapsing them into one score. The `final_ler_mean` rows test whether the affine surface reduces the protocol-defined error proxy under drift; the residual-boundary surrogate prevents that proxy from being overstated as finite-energy channel fidelity; the operation-count and fixed-point rows test whether the online rule is small enough for a real-time datapath; and the calibration rows test whether more than one estimator can populate the same committed interface. The joint claim is an interface-level one: adaptation, cost and traceability improve together only because the fast path remains a low-dimensional affine correction.

A successful FPGA experiment would therefore strengthen a specific physical interface claim, rather than simply add a hardware label to the present software result. The decisive measurements would be whether the same committed (K, b) banks reproduce the source-level affine corrections bit-accurately, whether the fast path meets the control-cycle latency budget with timing headroom, and whether the resource and power cost remain small compared with a per-shot neural or posterior-branch decoder. If those conditions are met, the main advantage over adjacent learned-decoder and real-time-decoder work would be a separation of roles: adaptation occurs in the slow calibration layer, where richer statistics or learned estimators can be used, whereas the latency-critical correction remains a deterministic physical displacement rule. Such a result would still not by itself prove a surface-code threshold, finite-energy logical-channel fidelity or broad holdout robustness, but it would turn the current affine-contract argument into a measured real-time calibration interface.

The scope of inference follows directly from this decomposition. The drift claim is strongest inside the predeclared four-scenario controlled simulation table and the two completed 12-pair scenario checks. The random-walk, burst/reset and faster-than-window stress tests broaden the diagnostic picture, but they use controlled known-state drift families rather than a trained-branch holdout benchmark. The nearest-syndrome, oracle-affine and wrapped-Gaussian checks separate estimator error, hard-correction risk and local affine validity, but they are not yet a full nearest-lattice or sequence-level known-noise benchmark. The residual-boundary surrogate gives channel-language bookkeeping for q/p half-lattice crossings, not finite-energy logical-channel fidelity.

The next evidence layer is therefore well defined. Statistically, the descriptive paired envelope should be replaced by repeat-expanded or inferential uncertainty estimates across all drift scenarios. Algorithmically, known-noise affine, nearest-lattice and sequence-level wrapped-Gaussian baselines

should be run under the same predeclared protocol. Physically, finite-energy logical-channel probabilities or channel infidelity should be estimated under specified noise channels before making fidelity-level claims. For hardware, the decisive evidence would be bit-accurate source-vs-board agreement, measured fast-path and commit latency, synthesis/resource reports and power measurements. Until those layers are added, the manuscript supports a controlled-simulation quantum-engineering systems method rather than a deployment result.

8 Conclusion

We presented a runtime-consistent affine calibration formulation for drift-adaptive GKP decoding. The method keeps the per-shot correction as the small affine rule $\Delta_t = K_t s_t + b_t$ while updating the affine parameters from recent syndrome statistics in a slower loop. Under a predeclared four-scenario controlled simulation protocol, the main-table hybrid residual branch has the lowest mean `final_ler_mean` in all tested scenarios. The current formal repeat-expanded evidence is narrower: 12-paired-repeat `static_bias_theta` and `linear_ramp` checks support positive scenario-level UKF-minus-hybrid intervals, while the remaining scenarios still require the same treatment before broad interval wording is justified. Scoped ablations and mechanism materials support a conservative interpretation: histogram dynamics and feature design matter, and the more general object is the affine calibration interface rather than teacher-parameter necessity. Controlled oracle-affine and wrapped-Gaussian checks show that known effective-state affine parameters reduce residual MSE in local-Gaussian states, while a simple wrapped posterior does not dominate in either one-step or short-sequence settings. A supplementary statistical-calibration analysis further supports the affine-contract view, while remaining separate from the main comparison.

Together, these results define a controlled-simulation systems method for drift-adaptive GKP correction. The contribution is the combination of a physical-layer calibration target, a deterministic low-cost fast path and a metric hierarchy that separates LER proxy, channel-language surrogate, software arithmetic and hardware-stage measurements. The distinctive advantage, relative to adjacent soft-information, learned-decoder and FPGA-decoder work, is the separation of adaptation from the latency-critical correction: the slow loop may use richer statistics or learned estimators, whereas the fast loop remains a small affine displacement rule that can be quantized, replayed and verified. Deployment-level and broad benchmark claims require the next evidence layer: measured device timing and resource use, default-runtime portability, expanded drift coverage, stronger baselines and repeat-expanded uncertainty estimates.

9 Outlook

Future validation should proceed along three axes. First, statistical benchmark expansion should add repeat-expanded uncertainty estimates, trained-branch holdout drift families and formal known-noise affine, nearest-lattice or sequence-level wrapped-Gaussian baselines under a predeclared protocol. Second, physics and channel validation should test the lattice constant, syndrome wrap interval, logical boundary and `final_ler_mean` definition before using finite-energy logical-channel simulations to support channel-level fidelity claims. Third, hardware validation should replace requirements with board data only after a future experiment records device paths, bitstream or RTL hashes, DMA or MMIO evidence, latency distributions, resource use, power and source-vs-board numerical agreement. Together, these studies would test whether the affine calibration contract remains advantageous beyond the present controlled simulation setting.

Supplementary reporting overview

The supplementary sections collect reporting details that make the manuscript traceable without expanding the empirical claims made in the Results and Discussion. They specify the data classes, protocol definitions, terminology, figure inputs and additional measurements needed before stronger statistical, channel-level or hardware conclusions could be drawn.

A Supplementary Literature Metric Crosswalk

The following tables record the representative external metrics used to orient the related-work comparison in the main text. They are comparison anchors, not normalized baselines or results generated by this study.

B Data, Code, and Reproducibility Availability

This appendix separates materials used by the present study from materials required for broader statistical, deployment or hardware claims. Available analysis artifacts support the controlled-simulation methods reported here, but they do not by themselves establish a hardware demonstration, a deployment study or a normalized broad-decoder leaderboard.

For a stronger empirical version of the study, the supporting material should include machine-readable result tables with seeds and uncertainty estimates, reproducible plotting scripts and figure manifests, exact configuration files for each drift family, runtime-artifact validation for the runner, a file-level supporting-file manifest for manuscript-facing tables and figures, a literature-metric crosswalk for related-work anchors, verified reference metadata and, when hardware language is used, board logs, bitstream hashes, fixed-point test vectors, latency measurements and resource reports. The present manuscript provides file-level hashes for its supporting CSV/JSON files and generator scripts, together with row metadata for the main controlled simulation benchmark rows. This remains narrower than a full independent reproduction package or a hardware traceability record. Until those stronger materials exist, hardware rows and broad reproducibility statements should be treated as implementation-stage validation targets.

C Scope of Reported Claims

D Benchmark Protocol and Supporting Tables

A formal numerical comparison must make clear whether results come from a fixed protocol or from post-hoc selection of convenient runs. Table 38 therefore separates the protocol, the supporting table and the manuscript use for each result class. This reporting table supports the controlled-simulation comparison; it does not turn the study into a hardware or deployment claim.

E Theory-to-Configuration Bridge

The theory sections use deliberately compact notation: a syndrome stream s_t , an affine fast-path update $\Delta_t = K_t s_t + b_t$, and a slow calibration loop that updates (K_t, b_t) from recent syndrome statistics. The present data instantiate this notation through four controlled synthetic drift scenarios in the predeclared simulation benchmark. These scenarios are protocol probes of an effective

calibration surface; they are not a complete physical noise model, a holdout drift distribution or a hardware environment.

The four scenario names should therefore be read as controlled synthetic probes. The static-bias scenario fixes an effective biased orientation state; the linear-ramp scenario tests monotone scale drift with orientation jitter; the step-change scenario tests an abrupt effective-state change; and the periodic-drift scenario tests oscillatory tracking. The protocol uses paired seeds, two repeats, four scenarios and five predeclared modes. This supports a scoped claim about runtime-consistent affine calibration under the predeclared simulation protocol; it does not support language about physical device drift, broad distributional generalization, confidence-level significance or real-board validation.

F Runtime Artifacts and Reproducibility Scope

The present study uses several artifact and runtime classes with different interpretations. Table 40 records the manuscript-facing taxonomy. This classification does not replace a full reproducibility package, but it states which artifacts support the reported simulation results and which artifacts remain outside the hardware evidence.

Three validation gaps remain. First, the present supporting-file checks cover the main result table, paired-delta tables, ablation and mechanism tables, statistical-calibration rows, controlled oracle/wrapped-Gaussian checks, sequence-controlled baselines, holdout stress diagnostics, fast-path cost, fixed-point parity, runtime-discipline counters, residual-boundary surrogate rows and the metric-readiness matrix. They still do not cover every narrative scope, availability or requirements table in the manuscript. Second, the supporting-file manifest records file-level SHA-256 hashes for manuscript-facing CSV/JSON files and generator scripts, and row metadata trace the main controlled simulation benchmark rows to scenario, mode, repeat, seed, source run path, config, runner and selected repeat-summary hashes. These records are still not a full independent reproduction package and do not hash every large event log. Third, the hardware surface remains a readiness check until a future board experiment produces device logs, bitstream or RTL hashes, DMA evidence, latency/resource measurements and source-vs-board vectors.

G Statistical Treatment and Reproducibility

The current main result is a predeclared-protocol ranking with descriptive uncertainty, not a distribution-level statistical claim. The main benchmark uses paired seeds and two repeats per scenario and mode. The manuscript therefore reports means, descriptive standard deviations, paired deltas and a descriptive paired-bootstrap envelope only. It does not report inferential confidence intervals, p -values, standard errors or hypothesis-test evidence.

The repeat-expansion route is therefore treated as a validation-threshold table rather than as additional evidence. Table 43 states which evidence class would be needed before the wording could move from descriptive ranking to repeat-expanded, holdout or hardware claims.

For the completed formal rows, the interval checks remain deliberately narrow. Table 44 reports the paired delta analysis for the completed `static.bias.theta` and `linear.ramp` rows. The lower bounds are positive for both scenarios, but the results do not replace the planned all-scenario paired-interval validation set, pooled analysis, holdout drift families or hardware-facing measurements.

These limits also determine the claim language in the abstract and conclusion. The supported statement is that the hybrid residual branch has the lowest descriptive mean `final.ler.mean` in the

predeclared controlled simulation set, and that the evidence motivates a runtime-consistent affine calibration strategy. The present evidence does not establish statistical significance, robustness across unseen drift families, full reproducibility, cross-host portability, hardware validation or deployment readiness.

H Claim-to-Evidence Summary

The manuscript is organized around a small number of claims. Each claim is tied to a specific result or supporting-material layer so that readers can separate method definitions, simulation results, diagnostic checks and hardware requirements. Tables 45 and 46 give the corresponding reporting route. They are supplementary summaries rather than additional experiments.

I Limitations and Required Validation

The validation limits are stated explicitly before they are used to qualify the claims. Table 47 summarizes the principal limitations of the present study and the evidence required before each claim could be strengthened.

J Terminology and Claim Boundaries

K Figure Data and Reporting Basis

The figure set is tied to underlying data tables and the stated claim scope. The present manuscript uses figure-generation scripts, and hardware validation is represented as a measurement-requirements table rather than as device data. The table below defines the intended claim and source-file basis for each figure in the present manuscript numbering.

References

- [1] Daniel Gottesman, Alexei Kitaev, and John Preskill. Encoding a qubit in an oscillator. *Physical Review A*, 64(1):012310, June 2001. arXiv:quant-ph/0008040.
- [2] Arne L. Grimsmo and Shruti Puri. Quantum Error Correction with the Gottesman-Kitaev-Preskill Code. *PRX Quantum*, 2(2):020101, June 2021. arXiv:2106.12989 [quant-ph].
- [3] Kosuke Fukui, Akihisa Tomita, Atsushi Okamoto, and Keisuke Fujii. High-threshold fault-tolerant quantum computation with analog quantum error correction. *Physical Review X*, 8(2):021054, 2018. eprint: 1712.00294.
- [4] Kyungjoo Noh and Christopher Chamberland. Fault-tolerant bosonic quantum error correction with the surface-GKP code. *Physical Review A*, 101(1):012316, 2020. eprint: 1908.03579.
- [5] Kyungjoo Noh, Christopher Chamberland, and Fernando G. S. L. Brandão. Low overhead fault-tolerant quantum error correction with the surface-GKP code. *PRX Quantum*, 3(1):010315, January 2022. arXiv:2103.06994 [quant-ph].
- [6] Nithin Raveendran, Narayanan Rengaswamy, Filip Rozpedek, Ankur Raina, Liang Jiang, and Bane Vasic. Finite Rate QLDPC-GKP Coding Scheme that Surpasses the CSS Hamming Bound. *Quantum*, 2022. eprint: 2111.07029.

- [7] Lucas Berent, Timo Hillmann, Jens Eisert, Robert Wille, and Joschka Roffe. Analog information decoding of bosonic quantum LDPC codes. *PRX Quantum*, 5(2):020349, May 2024. arXiv:2311.01328 [quant-ph].
- [8] Shantom K. Borah, Asit K. Pradhan, Nithin Raveendran, Michele Pacenti, and Bane Vasic. Fault Tolerant Decoding of QLDPC-GKP Codes with Circuit Level Soft Information, May 2025. arXiv:2505.06385 [quant-ph].
- [9] Kwok Ho Wan, Alex Neville, and W. S. Kolthammer. Memory-assisted decoder for approximate Gottesman-Kitaev-Preskill codes. *Physical Review Research*, 2(4):043280, November 2020. arXiv:1912.00829 [quant-ph].
- [10] Marc-Antoine Roy, Thomas Pousset, and Baptiste Royer. Decoding Multimode Gottesman-Kitaev-Preskill Codes with Noisy Auxiliary States, October 2025. arXiv:2510.12677 [quant-ph].
- [11] S. T. Spitz, B. Tarasinski, C. W. J. Beenakker, and T. E. O’Brien. Adaptive weight estimator for quantum error correction. *Advanced Quantum Technologies*, 2018. eprint: 1712.02360.
- [12] Thomas Wagner, Hermann Kampermann, Dagmar Bruss, and Martin Kliesch. Optimal noise estimation from syndrome statistics of quantum codes. *Physical Review Research*, 2021. eprint: 2010.02243.
- [13] Hanrui Wang, Pengyu Liu, Yilian Liu, Jiaqi Gu, Jonathan Baker, Frederic T. Chong, and Song Han. DGR: Tackling Drifted and Correlated Noise in Quantum Error Correction via Decoding Graph Re-weighting, 2023. eprint: 2311.16214.
- [14] Edward H. Chen, Theodore J. Yoder, Youngseok Kim, Neereja Sundaresan, Srikanth Srinivasan, Muyuan Li, Antonio D. Corcoles, Andrew W. Cross, and Maika Takita. Calibrated decoders for experimental quantum error correction. *Physical Review Letters*, 2022. eprint: 2110.04285.
- [15] Volodymyr Sivak, Michael Newman, and Paul Klimov. Optimization of decoder priors for accurate quantum error correction. *Physical Review Letters*, 133(15):150603, October 2024. arXiv:2406.02700 [quant-ph].
- [16] Johannes Bausch, Andrew W. Senior, Francisco J. H. Heras, Thomas Edlich, Alex Davies, Michael Newman, Cody Jones, Kevin Satzinger, Murphy Yuezhen Niu, Sam Blackwell, George Holland, Dvir Kafri, Juan Atalaya, Craig Gidney, Demis Hassabis, Sergio Boixo, Hartmut Neven, and Pushmeet Kohli. Learning high-accuracy error decoding for quantum processors. *Nature*, 635(8040):834–840, November 2024.
- [17] Christopher Chamberland, Jan Olle, Muyuan Li, Scott Thornton, and Igor Baratta. Fast and accurate AI-based pre-decoders for surface codes, 2026. eprint: 2604.12841.
- [18] Samuel Stein, Shuwen Kan, Chenxu Liu, Adrian Harkness, Sean Garner, Zefan Du, Yufei Ding, Ying Mao, and Ang Li. Calibration-Conditioned FiLM Decoders for Low-Latency Decoding of Quantum Error Correction Evaluated on IBM Repetition-Code Experiments, January 2026. arXiv:2601.16123 [quant-ph].
- [19] Haowen Wang, Yunjia Xue, Yingjie Qu, Xiaoyi Mu, and Hongyang Ma. Multidimensional Bose quantum error correction based on neural network decoder. *npj Quantum Information*, 8(1):134, November 2022.

- [20] Yotam Peled, David Zenati, and Eliya Nachmani. Neural Minimum Weight Perfect Matching for Quantum Error Codes, January 2026. arXiv:2601.00242 [quant-ph].
- [21] Maurice D. Hanisch, Bence Hetenyi, and James R. Wootton. Soft information decoding with superconducting qubits. *APS Open Science*, 1:000019, May 2026.
- [22] Volodymyr Sivak et al. Reinforcement Learning Control of Quantum Error Correction, March 2026. arXiv:2511.08493 [quant-ph].
- [23] Xiaohan Yang, Xuandong Sun, Zhiyi Wu, Jiawei Zhang, Ji Jiang, Xiayu Linpeng, Yuxuan Zhou, Ji Chu, Jingjing Niu, Youpeng Zhong, Song Liu, and Dapeng Yu. Real-time Surface-Code Error Correction Using an FPGA-based Neural-Network Decoder, May 2026. arXiv:2605.04892 [quant-ph] version: 1.
- [24] Poulami Das, Aditya Locharla, and Cody Jones. LILLIPUT: A Lightweight Low-Latency Lookup-Table Based Decoder for Near-term Quantum Error Correction, 2021. _eprint: 2108.06569.
- [25] Namitha Liyanage, Yue Wu, Alexander Deters, and Lin Zhong. Scalable Quantum Error Correction for Surface Codes using FPGA, 2023. _eprint: 2301.08419.
- [26] Ben Barber, Kenton M. Barnes, Tomasz Bialas, Okan Bugdayci, Earl T. Campbell, Neil I. Gillespie, Kauser Johar, Ram Rajan, Adam W. Richardson, Luka Skoric, Canberk Topal, Mark L. Turner, and Abbas B. Ziad. A real-time, scalable, fast and highly resource efficient decoder for a quantum computer, 2023. _eprint: 2309.05558.
- [27] Satvik Maurya, Thilo Maurer, Markus Buhler, Drew Vandeth, and Michael E. Beverland. FPGA-tailored algorithms for real-time decoding of quantum LDPC codes, 2025. _eprint: 2511.21660.
- [28] Laura Caune, Luka Skoric, Nick S. Blunt, Archibald Ruban, Jimmy McDaniel, Joseph A. Valery, Andrew D. Patterson, Alexander V. Gramolin, Joonas Majaniemi, Kenton M. Barnes, Tomasz Bialas, Okan Buğdaycı, Ophelia Crawford, György P. Gehér, Hari Krovi, Elisha Matekole, Canberk Topal, Stefano Poletto, Michael Bryant, Kalan Snyder, Neil I. Gillespie, Glenn Jones, Kauser Johar, Earl T. Campbell, and Alexander D. Hill. Demonstrating real-time and low-latency quantum error correction with superconducting qubits, October 2024. arXiv:2410.05202 [quant-ph].
- [29] Abbas B. Ziad, Ankit Zalawadiya, Canberk Topal, Joan Camps, Gyorgy P. Geher, Matthew P. Stafford, and Mark L. Turner. Local Clustering Decoder: a fast and adaptive hardware decoder for the surface code, 2024. _eprint: 2411.10343.
- [30] Thilo Maurer, Markus Buhler, Michael Kroner, Frank Haverkamp, Tristan Muller, Drew Vandeth, and Blake R. Johnson. Real-time decoding of the gross code memory with FPGAs, 2025. _eprint: 2510.21600.
- [31] Mahnaz Jafarzadeh, Jonathan Conrad, Rafael N. Alexander, and Ben Q. Baragiola. Logical channels in approximate Gottesman-Kitaev-Preskill error correction. *Physical Review A*, 112(6):062413, December 2025. arXiv:2504.13383 [quant-ph].
- [32] Jacob Hastrup and Ulrik L. Andersen. Analysis of loss correction with the Gottesman-Kitaev-Preskill code. *Physical Review A*, 108(5):052413, November 2023. arXiv:2112.01425 [quant-ph].

- [33] Guo Zheng, Wenhao He, Gideon Lee, Kyungjoo Noh, and Liang Jiang. Performance and achievable rates of the Gottesman-Kitaev-Preskill code for pure-loss and amplification channels, December 2024. arXiv:2412.06715 [quant-ph].

Prior work class	Reported LER, fidelity, latency or resource metric	This work reports	Not yet measured here
Analog and surface-GKP decoding	Squeezing, failure and overhead anchors include about $16.0 \rightarrow 9.8$ dB required squeezing, $11.2/18.6$ dB and $0.69/0.81\%$ threshold-type values, $0.87\% \rightarrow 0.36\%$ CNOT-failure reduction, and 291 modes plus 97 qubits versus 1457 qubits for a matched target [3, 4, 5, 7, 8].	A physical-layer affine calibration rule driven by recent GKP syndrome histograms and evaluated by a protocol-defined residual-boundary LER proxy.	End-to-end surface-GKP threshold, overhead or logical-channel failure rate.
Calibration-aware and learned QEC decoders	Drift-aware graph and prior calibration report $3.6\times$ average mismatch-LER reduction, worst-case reductions up to $7360\times$, and prior-calibration gains of 48%, 16% and 3.3%; calibration-conditioned learned decoders report large-shot real-device studies and up to $11.11\times$ LER reductions [13, 14, 15, 16, 18].	A learned slow-loop estimator for a low-dimensional affine surface; the per-shot GKP correction remains deterministic.	Real-processor learned-decoder LER, calibration-shot scaling or device-dataset generalization.
Finite-energy logical-channel analyses	Logical-channel probabilities, channel infidelity and loss/amplification recovery curves include anchors such as GRN Pauli probabilities $p_I = 0.9900$, logical-channel infidelity curves and near-optimal loss/amplification comparisons [31, 32, 33].	Residual-boundary event decompositions and a finite-squeezing toy-channel sanity check that connect the LER proxy to channel language.	Calibrated finite-energy logical-channel fidelity, channel infidelity or process tomography.
Runtime pre-decoders and calibration-conditioned neural modules	Runtime anchors include $270.83 \rightarrow 38.78 \mu\text{s}$ sub-step pre-decoding, $85\text{--}95 \mu\text{s}$ folded FiLM latency, and real-time neural feedback examples with 124 ns neural decoding inside a 550 ns closed loop [17, 18, 23].	An analytical fast path with four multiplications, four additions, no nonlinear operation and Q4.20 software parity in controlled samples.	Per-shot neural-decoder latency, closed-loop feedback latency or on-device FPGA timing.
Real-time FPGA and hardware-tailored decoders	Hardware papers report 28–42 ns lookup latency, 11.5 ns/round FPGA decoding, 810/240 ns FPGA/ASIC collision decoding, sub- μs feedback paths, memory compression, area, power and source-to-device evidence [24, 25, 27, 28, 29, 30, 23].	A compact, fixed-point-checkable affine target with software runtime counters and explicit board-measurement requirements.	FPGA synthesis, LUT/FF/DSP/BRAM use, power, real-board latency and source-vs-board agreement.

Table 2: Closest-work metric comparison. The table separates what adjacent work reports, what this manuscript reports and which validation class remains unmeasured. The comparison is intentionally not a normalized leaderboard, because the code families, physical endpoints and hardware-validation levels differ.

Metric axis	Manuscript metric	Supported statement	Missing for stronger claim
Logical-error proxy	<code>final_1er_mean</code> means, paired descriptive deltas and a non-inferential paired-bootstrap envelope	Hybrid residual calibration has the lowest mean protocol-defined <code>final_1er_mean</code> in the tested drift scenarios.	More repeats, inferential confidence intervals, holdout drift and logical-channel reconstruction.
Logical-channel fidelity	Residual-boundary Pauli-channel-style surrogate, method-level lattice sanity summary and finite-squeezing toy-channel sanity check	The surrogate links q/p half-lattice crossing rates to channel language, while the toy finite-squeezing check compares hard nearest-syndrome, fixed affine and oracle affine rules under a simplified noisy measurement channel; neither estimates finite-energy logical-channel fidelity or process infidelity.	Channel-specific calibrated finite-energy simulation or logical-channel tomography.
Drift robustness	Four scenario rankings plus controlled local-Gaussian checks, holdout drift stress tests and a commit-lag sweep	Controlled stress diagnostics expose where stale commits help or fail under random-walk, burst/reset and faster-than-window settings; they do not establish trained-branch holdout robustness.	A formal predeclared expanded benchmark, repeat-expanded uncertainty estimates, trained-branch holdout validation and board-measured stale-parameter latency distributions.
Fast-path cost and latency	Analytical per-shot operation count plus Q4.20 software parity	The affine correction has lower analytical arithmetic and state footprint than the wrapped-posterior references evaluated here, and fixed-point emulation introduces no residual-boundary crossing change in the controlled sample set.	FPGA synthesis, timing closure, resource and power measurements, source-vs-board agreement and board-level latency distribution.
Hardware-facing validation	Measurement requirements only	The manuscript defines what must be measured before making hardware claims.	Board logs, bitstream or RTL hash, DMA evidence, commit reliability and source-vs-board vectors.

Table 3: Metric-readiness matrix for the manuscript. The table separates reported metrics from external standards and implementation-stage validation requirements; it does not convert the `final_1er_mean` proxy into channel fidelity, statistical significance or hardware latency.

Quantity	Meaning in the manuscript	Not claimed here
Conventional GKP spacing	Used to introduce the ideal square-lattice code and finite-energy comb picture.	A calibrated device-specific quadrature scale.
$\lambda = \sqrt{2\pi}$	Software residual wrap period used by the emulator, tables and supporting-data checks.	A new physical code definition.
$\lambda/2$ boundary	Normalized q/p half-cell boundary for the event indicator in Eq. (3).	Outer-code logical failure probability.
<code>final_ler_mean</code>	Empirical mean of residual half-cell crossings under the stated runtime-calibration protocol.	Finite-energy logical-channel fidelity, process fidelity or tomography.

Table 4: Coordinate and metric convention. The table separates the physical GKP motivation from the normalized software residual coordinate used in the reported runtime-calibration benchmark.

σ_r	Eq. squeezing (dB)	p_{cross}	p_{any}	Surrogate infidelity
0.15	16.48	0.000000	0.000000	0.000000
0.20	13.98	0.000000	0.000000	0.000000
0.25	12.04	0.000001	0.000001	0.000001
0.30	10.46	0.000029	0.000059	0.000039
0.35	9.12	0.000342	0.000685	0.000456
0.40	7.96	0.001729	0.003454	0.002303

Table 5: Analytical residual-boundary sensitivity for a zero-mean Gaussian residual model. The equivalent squeezing column is $-10 \log_{10} \sigma_r^2$. The surrogate infidelity is $1 - (1 + 2(1 - p_{\text{any}}))/3$, used only to connect boundary events to channel-style terminology; it is not a finite-energy GKP logical-channel fidelity or process-infidelity estimate.

Object	Current support	Manuscript use	Missing before stronger claim
Syndrome wrapping	Shared $\sqrt{2\pi}$ lattice constant and wrap interval $[-\lambda/2, \lambda/2)$.	Approximate GKP syndrome model.	Independent symbol-to-code validation.
Finite-energy and measurement noise	Configurable Δ , measurement efficiency, shot noise and ancilla-error terms in the software measurement model.	Approximate measurement effects only.	Full physical-noise calibration.
Effective drift	Synthetic $(\sigma, \mu_q, \mu_p, \vartheta)$ trajectories drive the controlled simulation protocol.	Controlled drift probes.	Unseen or device-derived drift families.
<code>final_ler_mean</code>	Residual accumulation is scored against half-lattice logical boundaries in the fast-loop emulator.	Predeclared-protocol logical-error proxy.	Logical-channel validation and stronger baselines.
Hardware path	Board-facing rows specify future measurement requirements.	Future validation target.	Measured device logs and resource/timing data.

Table 6: Physics-model and metric scope. The current implementation anchors the notation used in the paper, but the evidence supports only a controlled simulation protocol claim.

Step	Operation	Reproducible state recorded by the protocol
1	Accumulate recent syndrome statistics.	Windowed histogram $H_{t-c+1:t}$, histogram differences, active scenario index and fast-loop shot count.
2	Propose a candidate affine surface.	Candidate K_t^{cand} , b_t^{cand} , estimator family and any teacher or statistical state used to compute the proposal.
3	Apply runtime guards before exposure to the fast path.	Clipped and optionally smoothed K, b , saturation flags, overflow flags and the inactive bank selected for staging.
4	Commit only at a safe update boundary.	Commit acknowledgement, active-bank selector, stale-parameter counter and rollback status.
5	Correct every shot with the committed affine rule.	Clipped syndrome, quantized correction, residual displacement and q/p half-lattice boundary indicators.

Table 7: Reproducible algorithmic contract for affine runtime calibration. Different slow-loop estimators may populate the candidate surface, but all paths must expose the same committed K, b state and the same fast-loop log fields. These fields are sufficient for the software comparisons reported here and define the future source-vs-board trace needed for hardware validation.

Datapath assumption	Parallel schedule	Serialized schedule
Numeric format	Q4.20 fixed-point words, matching the software parity check.	Same Q4.20 source-vector format.
Multiplier allocation	Four fixed-point multipliers for the 2×2 product.	One or two multipliers reused across the four products.
Cycle model	Product stage followed by add, bias, clip and register stages; clock period not claimed.	More arithmetic cycles, with the same input and output tuple stream; clock period not claimed.
Active state	Six affine words (K, b) plus clip thresholds, bank selector and status flags.	Same state; lower multiplier footprint changes schedule, not protocol state.
Validation target	Match source corrections and flags within declared fixed-point tolerance.	Same source-vs-board target.

Table 8: Pre-synthesis datapath assumptions for the affine fast path. The table states the implementation model that motivates the operation-count and fixed-point checks. It is not synthesis, timing closure, resource measurement or a board result.

Datapath element	Manuscript contract	Supporting analysis
Input state	Wrapped syndrome s_t , clip limits and active-bank selector.	Defined by the software fast-loop interface and metric definition.
Parameter bank	Active and inactive K, b banks with staged commit and rollback events.	Exercised by controlled simulation runtime counters only.
Arithmetic	One 2×2 matrix-vector operation, one bias vector and bounded clipping in fixed-point form.	Analytical operation count and Q4.20 software fixed-point emulation.
Output check	Corrected displacement and residual-boundary crossing indicator.	Compared in software against the floating-point reference.
Device evidence	Timing, resource, power and source-vs-board agreement.	Outside the present study; required for a hardware claim.

Table 9: FPGA-facing datapath contract for the affine fast path. The table defines what a future board implementation would need to reproduce. The interface is supported here by software and analytical checks only; the study does not report synthesis, timing closure or board measurements.

Item	Reported setting	Scope
Input tensor	Five 32×32 histogram windows, four histogram-delta planes, teacher state, teacher runtime bias and teacher temporal deltas, encoded as 21 spatial channels.	Slow-loop residual branch only.
Targets	Two residual- b labels, b_q and b_p , defined as target runtime bias minus teacher runtime bias.	Residual calibration, not direct logical decoding.
Dataset split	2048 runtime windows from the four drift scenarios, split 1638/204/206 into train/validation/test windows.	Split-level check only; not unseen-drift validation.
Model	One 3×3 convolution with 16 channels, ReLU, 2×2 average pooling, one 96-unit hidden layer and two outputs.	Reported float model architecture.
Training	Weighted MSE on normalized residual- b outputs, label weights $w_{b_q} = w_{b_p} = 2$, learning rate 2×10^{-3} , weight decay 10^{-4} , batch size 64, 36-epoch cap and patience 8.	Training description only; not full reproducibility proof.

Table 10: Method details for the teacher-anchored residual branch. The table documents the reported slow-loop CNN component used by the manuscript; it does not make the branch a per-shot neural decoder or establish holdout-drift generalization.

Comparator or diagnostic	Question addressed	Interpretation limit
EKF and UKF modes	Whether standard nonlinear filtering baselines are competitive under the same scenario matrix and metric.	They are not optimized full posterior GKP decoders or hardware implementations.
Constant- μ and RLS residual- b modes	Whether simple affine or online residual updates explain the result without the teacher-anchored branch.	They do not cover all possible adaptive-prior or known-noise estimators.
Teacher-anchored hybrid residual- b	Whether a slow-loop residual branch can improve the committed affine fast path in the predeclared protocol.	Two repeats per scenario support only descriptive ranking, not inferential significance.
Oracle affine checks	Whether the affine mapper has headroom when the effective drift state is known.	Oracle state access is a diagnostic upper reference, not a deployable estimator.
Wrapped-Gaussian mean and MAP checks	Whether a simple branch-aware posterior rule obviously dominates the affine correction in local controlled states.	One-step or short-sequence controlled references do not replace a formal sequence-level wrapped-decoder benchmark.
Holdout drift stress families	Whether random-walk, burst/reset and faster-than-window drift expose adaptation or stale-commit failures.	The stress rows use controlled known-state references, not trained-branch holdout generalization.
Supplementary statistical calibration	Whether the same affine interface can be populated by a non-neural estimator.	It supports estimator interchangeability, not qualification of statistical calibration as a validated main comparator.

Table 11: Comparator ladder for the simulation and diagnostic evidence. The table states what each comparator can and cannot establish, so that the reported `final_ler_mean`, residual-boundary and cost results are read at the proper validation level.

Scenario	EKF	UKF	Const.- μ	RLS- b	Hybrid- b
<code>static.bias.theta</code>	0.838110	0.825370	0.836658	0.837577	0.810902
<code>linear.ramp</code>	0.819200	0.811201	0.816911	0.819373	0.787755
<code>step.sigma.theta</code>	0.822365	0.811548	0.819784	0.821493	0.788800
<code>periodic.drift</code>	0.832192	0.821558	0.829670	0.832334	0.806392

Table 12: Predeclared four-scenario controlled simulation benchmark. Entries are `final_ler_mean`; lower is better. The table supports a predeclared-protocol controlled simulation ranking. Confidence intervals, formal hypothesis tests and a predeclared expanded holdout benchmark remain open validation requirements.

Scenario	Mean Δ final_1er_mean vs UKF	Relative reduction	Minimum paired Δ
static_bias_theta	0.014469	1.75%	0.013953
linear_ramp	0.023446	2.89%	0.023017
step_sigma_theta	0.022748	2.80%	0.022440
periodic_drift	0.015166	1.85%	0.013571

Table 13: Paired descriptive UKF-vs-hybrid deltas from the two available repeats per scenario. Δ final_1er_mean is UKF minus Hybrid-*b*, so positive values indicate lower final_1er_mean for the hybrid branch. The table documents repeat-level directional consistency in the existing supporting data; it is not a confidence interval, *p*-value or distribution-level robustness test.

Scenario	<i>n</i> pairs	Mean Δ	Envelope low	Envelope high	Positive pairs
static_bias_theta	2	0.014469	0.013953	0.014984	2/2
linear_ramp	2	0.023446	0.023017	0.023875	2/2
step_sigma_theta	2	0.022748	0.022440	0.023056	2/2
periodic_drift	2	0.015166	0.013571	0.016762	2/2
all_scenarios	8	0.018957	0.016044	0.021892	8/8

Table 14: Descriptive paired-delta resampling envelope for the UKF-versus-hybrid comparison. Δ is UKF final_1er_mean minus Hybrid-*b* final_1er_mean for paired repeats, so positive values indicate lower hybrid final_1er_mean. The envelope is a non-inferential paired-bootstrap percentile summary over the available repeat-level deltas. With $n = 2$ per scenario it is reported only as a descriptive summary and must not be read as a confidence interval, *p*-value, standard error or robustness claim.

Scenario	UKF mean	Hybrid mean	Mean Δ	Relative reduction	Positive pairs	Δ /max SD
static_bias_theta	0.825370	0.810902	0.014469	1.75%	2/2	12.17
linear_ramp	0.811201	0.787755	0.023446	2.89%	2/2	27.00
step_sigma_theta	0.811548	0.788800	0.022748	2.80%	2/2	21.28
periodic_drift	0.821558	0.806392	0.015166	1.85%	2/2	8.05

Table 15: Supporting-data-backed descriptive UKF-minus-Hybrid advantage margins. Δ is UKF final_1er_mean minus Hybrid-*b* final_1er_mean; positive values therefore indicate lower hybrid final_1er_mean. The last column divides the mean paired margin by the larger reported descriptive SD of the UKF and Hybrid-*b* rows in the same scenario. This is an effect-size-style scale readout for diagnostic readability only and is not a confidence interval, standard error, *p*-value, significance test, expanded benchmark or hardware measurement.

Scenario	UKF <code>final_ler_mean</code>	Hybrid <code>final_ler_mean</code>	Δ	Relative reduction	Positive pairs
<code>static_bias_theta</code>	0.817940	0.801012	0.016927	2.07%	2/2
<code>linear_ramp</code>	0.840509	0.813909	0.026600	3.16%	2/2
<code>step_sigma_theta</code>	0.828792	0.816023	0.012769	1.54%	2/2
<code>periodic_drift</code>	0.810660	0.781844	0.028817	3.55%	2/2

Table 16: Short-run protocol check for the planned repeat-expanded anchor comparison. Δ is UKF `final_ler_mean` minus Hybrid-*b* `final_ler_mean`, so positive values indicate lower hybrid `final_ler_mean`. The source run used a shortened configuration (120 slow-loop updates and 4.8×10^5 fast cycles), two paired repeats and the formal runner over all four predeclared scenarios. This table checks execution coverage and field availability for the expansion route only; it is not the main benchmark, not an expanded benchmark, not an inferential interval and not hardware evidence.

Controlled state	Nearest syndrome	Fixed affine	Oracle affine	Wrapped mean	Wrapped MAP
<code>static_bias_theta</code>	0.232823	0.049199	0.048095	0.047955	0.048203
<code>linear_ramp_midpoint</code>	0.230689	0.061417	0.060251	0.062341	0.063105
<code>step_after_jump</code>	0.233626	0.108216	0.092541	0.105815	0.110339
<code>periodic_high_phase</code>	0.231307	0.081757	0.076028	0.082349	0.084002

Table 17: Controlled local-Gaussian nearest-syndrome, oracle-affine and wrapped-Gaussian checks. Entries are one-step residual MSE; lower is better. The nearest-syndrome row applies the wrapped syndrome directly as a hard correction and is included as a one-step nearest-lattice-style sanity reference, not as a tuned finite-energy nearest-lattice decoder. The table uses the manuscript affine mapper for the affine rows and known-state wrapped-Gaussian posterior rules for the posterior rows. It is a sanity check for estimator and decoder limits, not a formal controlled simulation benchmark, sequence-level wrapped decoder comparison, holdout-drift test or hardware measurement.

Controlled state	Fixed affine	Oracle affine	Wrapped mean	Wrapped MAP
<code>static_bias_theta</code>	0.000000	0.000000	0.000000	0.000017
<code>linear_ramp_midpoint</code>	0.000000	0.000000	0.000042	0.000183
<code>step_after_jump</code>	0.000467	0.000467	0.002417	0.006450
<code>periodic_high_phase</code>	0.000042	0.000042	0.000617	0.001867

Table 18: Residual-boundary Pauli-channel-style surrogate for the controlled local-Gaussian states. Entries are p_{any} , the probability of at least one q/p half-lattice residual crossing; lower is better. The source CSV also records q-only, p-only, simultaneous crossing and $F_{\text{avg}}^{\text{surr}}$ fields, but the table reports only the crossing probability to avoid presenting the surrogate as finite-energy logical-channel tomography or process fidelity.

Controlled state	Fixed affine	Oracle affine	Wrapped mean	Wrapped MAP
<code>static_bias_theta</code>	1.000000	1.000000	1.000000	0.999989
<code>linear_ramp_midpoint</code>	1.000000	1.000000	0.999972	0.999878
<code>step_after_jump</code>	0.999689	0.999689	0.998389	0.995700
<code>periodic_high_phase</code>	0.999972	0.999972	0.999589	0.998756

Table 19: Pauli-style identity-surrogate readout derived from the same residual-boundary event decomposition as Table 18. Entries are $F_{\text{avg}}^{\text{surr}} = (1 + 2p_I^{\text{surr}})/3$, with $p_I^{\text{surr}} = 1 - p_{\text{any}}$. The table is included only to make the manuscript’s channel-language bridge traceable to supporting data; it is not a finite-energy GKP logical-channel fidelity, process fidelity, hardware fidelity or outer-code logical-error estimate.

Method	Mean p_{any}	Worst state	Worst p_{any}	Mean $F_{\text{avg}}^{\text{surr}}$	Worst $F_{\text{avg}}^{\text{surr}}$
Fixed affine	0.000127	<code>step_after_jump</code>	0.000467	0.999915	0.999689
Oracle affine	0.000127	<code>step_after_jump</code>	0.000467	0.999915	0.999689
Wrapped mean	0.000769	<code>step_after_jump</code>	0.002417	0.999488	0.998389
Wrapped MAP	0.002129	<code>step_after_jump</code>	0.006450	0.998581	0.995700

Table 20: Method-level lattice sanity summary for the residual-boundary channel surrogate. Values aggregate the same controlled local-Gaussian states and source rows as Tables 18 and 19. The table is included to make the channel-language diagnostic easier to check; it is not finite-energy GKP logical-channel fidelity, process tomography, hardware fidelity or an outer-code logical-error estimate.

Δ	Method	Mean p_{any}	Worst state	Worst p_{any}	Mean $F_{\text{avg}}^{\text{surr}}$
0.18	Hard nearest-syndrome	0.000148	<code>step_after_jump</code>	0.000517	0.999901
0.18	Fixed affine	0.000148	<code>step_after_jump</code>	0.000517	0.999901
0.18	Oracle affine	0.000148	<code>step_after_jump</code>	0.000517	0.999901
0.26	Hard nearest-syndrome	0.000160	<code>step_after_jump</code>	0.000475	0.999893
0.26	Fixed affine	0.000129	<code>step_after_jump</code>	0.000458	0.999914
0.26	Oracle affine	0.000129	<code>step_after_jump</code>	0.000458	0.999914
0.34	Hard nearest-syndrome	0.001002	<code>step_after_jump</code>	0.001500	0.999332
0.34	Fixed affine	0.000163	<code>step_after_jump</code>	0.000617	0.999892
0.34	Oracle affine	0.000165	<code>step_after_jump</code>	0.000625	0.999890

Table 21: Finite-squeezing toy-channel sanity check. The table adds a scalar finite-squeezing width Δ to a simplified noisy wrapped-syndrome measurement channel and reports residual half-lattice event probabilities for three correction rules. It is a toy measurement-channel consistency check only, not a calibrated finite-energy GKP logical-channel fidelity estimate, process tomography, formal benchmark, hardware fidelity or FPGA validation.

Scenario	Nearest syndrome	Fixed affine	Oracle affine	Wrapped mean	Wrapped MAP
static.bias.theta	0.000214	0.000000	0.000000	0.000000	0.000000
linear.ramp	0.000264	0.000015	0.000015	0.000158	0.000575
step.sigma.theta	0.000580	0.000310	0.000310	0.001221	0.003220
periodic.drift	0.000229	0.000005	0.000005	0.000127	0.000519

Table 22: Sequence-level controlled baseline analysis. Entries are mean half-lattice residual-boundary crossing proxies across 384 sequences of 512 steps; lower is better. The nearest-syndrome row is a direct hard-correction reference rather than a tuned nearest-lattice decoder. This is a controlled local-Gaussian sequence analysis, not the formal controlled simulation benchmark, a confidence-interval run, a holdout drift test or hardware validation.

Holdout stress family	Fixed affine	Lagged affine	Oracle affine	Wrapped mean	Wrapped MAP	Oracle $F_{\text{avg}}^{\text{sur}}$
random.walk.drift	0.078869	0.073867	0.072685	0.078740	0.080857	0.999915
burst.reset.drift	0.066865	0.068622	0.062144	0.065936	0.067789	0.999780
faster.than.window.oscillation	0.068354	0.068890	0.063745	0.067485	0.068764	0.999969

Table 23: Controlled known-state holdout drift stress diagnostic. Entries are residual MSE across 384 sequences of 512 steps unless otherwise stated; lower residual MSE is better. $F_{\text{avg}}^{\text{sur}}$ is the Pauli-channel-style surrogate induced by the oracle-affine residual-boundary identity rate, not a finite-energy logical-channel fidelity. The lagged-affine row uses known-state affine parameters committed every 64 steps with a 32-step lag, and is a slow-loop lag reference rather than the trained residual branch. The table is a non-hardware stress test for the affine contract, not a formal expanded benchmark, confidence-interval analysis or process-fidelity estimate.

Surface	Layer	Oracle MSE gain (%)	Branch-risk delta	Lag-risk delta	Readout
static.bias.theta	Short sequence	2.05	0.000000	–	local affine not dominated
linear.ramp	Short sequence	3.40	0.000142	–	local affine headroom visible
step.sigma.theta	Short sequence	10.00	0.000910	–	local affine headroom visible
periodic.drift	Short sequence	3.43	0.000122	–	local affine headroom visible
random.walk.drift	Holdout stress	7.84	0.006055	0.001182	local affine headroom visible
burst.reset.drift	Holdout stress	7.06	0.003792	0.006478	stale commit can erase gain
faster.than.window.oscillation	Holdout stress	6.74	0.003740	0.005144	stale commit can erase gain

Table 24: Affine local-validity diagnostic derived from the controlled sequence and holdout-stress CSVs. Oracle MSE gain is $(\text{MSE}_{\text{fixed}} - \text{MSE}_{\text{oracle}})/\text{MSE}_{\text{fixed}}$. For short-sequence rows, branch-risk delta is the wrapped-posterior-mean crossing proxy minus the oracle-affine crossing proxy. For holdout-stress rows, branch-risk and lag-risk deltas are residual-MSE differences relative to oracle affine. The table is a derived diagnostic for the local affine approximation, not a formal nearest-lattice, known-noise or wrapped-decoder benchmark, not a confidence-interval analysis and not hardware evidence.

Holdout stress family	Lag 0	Lag 8	Lag 16	Lag 32	Lag 64	Lag 128
random.walk.drift	0.073831	0.073976	0.074130	0.074410	0.075173	0.076093
burst.reset.drift	0.067454	0.068193	0.068657	0.068694	0.068583	0.066913
faster.than.window.oscillation	0.069251	0.069406	0.068989	0.069039	0.068842	0.068249

Table 25: Commit-lag sweep for the lagged affine reference. Entries are residual MSE across 384 sequences of 512 steps; lower is better. The commit interval is fixed at 64 simulation steps and the columns sweep additional simulation-step lag. This table is a non-hardware stale-parameter diagnostic, not a measured FPGA-latency result, source-vs-board agreement, trained-branch holdout generalization, a confidence-interval analysis or logical-channel fidelity.

Decoder path	Branches	Mult.	Add.	Nonlinear ops	State scalars
Affine fast path	1	4	4	0	6
Wrapped MAP, 3x3 branches	9	49	40	0	33
Wrapped posterior mean, 3x3 branches	9	99	98	18	33

Table 26: Analytical per-shot operation-count model. The affine row counts the two-quadrature matrix-vector correction used in the manuscript fast path. The wrapped rows count one-step known-state 3-by-3 branch references. The table supports a low-complexity design rationale only; it is not FPGA synthesis, timing closure, power/resource measurement or hardware validation.

Scenario	Max diff	p99 diff	MSE delta	Crossing delta	Quant. sat.
static_bias_theta	0.000001	0.000001	0.000000	0.000000	0.000000
linear_ramp_midpoint	0.000001	0.000001	0.000000	0.000000	0.000000
step_after_jump	0.000002	0.000001	0.000000	0.000000	0.000000
periodic_high_phase	0.000001	0.000001	0.000000	0.000000	0.000000

Table 27: Software fixed-point parity check for the affine fast path. Entries compare Q4.20 fixed-point emulation against the floating-point runtime under a controlled local-Gaussian sample set. Max and p99 diff are absolute correction differences. MSE delta and crossing delta are fixed-point minus floating-point. This table is a numerical-emulation check only, not FPGA synthesis, timing closure, source-vs-board agreement or hardware validation.

Mode	Commits	Slow viol.	Fast viol.	Overflow	Corr. sat.
Constant Residual-Mu	899.8	0	0.0000158	0.002582	0
EKF	899.8	0	0.0000158	0.002559	0
Hybrid Residual-B	899.9	0	0.0000158	0.002536	0
RLS Residual-B	899.8	0	0.0000158	0.002538	0
UKF	899.8	0	0.0000158	0.002571	0

Table 28: Software runtime-discipline counters averaged across scenario-level comparison rows. The table reports software-in-the-loop commit counts, slow-update violation rate, fast-cycle violation rate, overflow rate and correction-saturation rate. It supports software observability of the stage-and-commit contract only; it is not board commit latency, hardware reliability or FPGA timing evidence.

Mode	Avg. LER	Δ vs UKF	Δ vs full hybrid
ukf	0.817382	0.000000	+0.018837
hybrid_full	0.798545	-0.018837	0.000000
hybrid_no_hist_deltas	0.826723	+0.009341	+0.028178
hybrid_no_teacher_prediction	0.807251	-0.010131	+0.008706
hybrid_no_teacher_params	0.749621	-0.067761	-0.048924
hybrid_no_teacher_deltas	0.800329	-0.017053	+0.001784

Table 29: Scoped feature and teacher ablation. Negative deltas indicate lower LER than the reference. These rows support feature sensitivity, not causal teacher necessity.

Seed	Gated v5 – Full	I1 – Gated v5	Verdict
20260425	0.000907	0.163289	harmful
20260427	-0.145352	0.287166	harmful
20260428	-0.078998	0.057395	harmful
20260429	-0.127948	0.322245	harmful
20260430	-0.170777	-0.024372	mixed/no clear effect
20260510	-0.003953	-0.035533	helpful

Table 30: Descriptive six-seed mechanism and intervention summary. The lower-clip intervention does not identify a causal mechanism and should not be described as a mitigation.

Scenario	UKF	Hybrid- <i>b</i>	StatCalib extension (not matched)
<code>static_bias_theta</code>	0.825370	0.810902	0.431708
<code>linear_ramp</code>	0.811201	0.787755	0.467083
<code>step_sigma_theta</code>	0.811548	0.788800	0.460016
<code>periodic_drift</code>	0.821558	0.806392	0.438751

Table 31: Supplementary statistical-calibration analysis. The table documents a separate estimator-protocol observation. Because candidate generation and selection are not matched to the main ranking protocol, these rows support the generality of the affine interface but do not establish a fully matched comparator, a unique clean threshold or a fair main-table ranking. The absolute values are not comparable to the main ranking because candidate selection and protocol generation differ.

Hardware-facing item	Manuscript treatment	Evidence needed for hardware claims
Board platform and host	<i>Not reported</i>	Linux + FPGA host, device path, permissions, board model and driver version.
Bitstream / RTL / DMA interface	<i>Not reported</i>	Bitstream hash, AXI map, DMA histogram shape, element width and timeout policy.
Fast-path latency	<i>Not reported</i>	p50/p95/p99 and worst-case cycles for $Ks + b$, clipping and bank selection.
Commit latency and reliability	<i>Not reported</i>	stage, commit, acknowledgement, rollback and stale-parameter counters.
Resource and power	<i>Not reported</i>	LUT/FF/DSP/BRAM, clock frequency, timing closure and power.
Source-vs-board agreement	<i>Not reported</i>	fixed-point agreement between software reference and board output on shared test vectors.

Table 32: Hardware-measurement requirements. These entries define the implementation-stage evidence required before hardware claims can be made; they are not results of the present study.

Advantage dimension	Supported basis in this study	Relation to prior work	Required upgrade
Latency-critical role	The per-shot rule is a 2×2 affine map plus clipping and bank selection; the learned estimator is outside the per-shot path.	Complements AI pre-decoders and calibration-conditioned learned decoders by placing learning in a slow calibration loop rather than in the correction datapath.	Measured board latency distribution.
Noise adaptation	The hybrid branch ranks lowest by mean <code>final_ler_mean</code> in the four tested drift scenarios; controlled stress tests expose random-walk, burst/reset and faster-than-window regimes.	Addresses a physical-layer GKP correction surface rather than only outer-code graph weights, decoder priors or message-passing reliabilities.	Repeat-expanded holdout benchmark.
Cost and verifiability	The affine count is 4 multiplications, 4 additions and 6 stored state scalars per shot, with Q4.20 software parity under controlled samples; parallel or serialized multiplier schedules preserve the same source trace.	Gives a smaller source-vs-board verification target than a branch-aware posterior rule or a per-shot neural decoder.	Synthesis, resource and power reports.
GKP metric interpretation	<code>final_ler_mean</code> counts q/p half-lattice boundary crossings, and the residual-boundary surrogate decomposes the same events into channel-language categories.	Connects the protocol to approximate-GKP logical-boundary reasoning while remaining narrower than finite-energy channel fidelity.	Logical-channel reconstruction.
Hardware extension path	The manuscript specifies timing, commit, resource, power and source-vs-board measurements needed for an FPGA demonstration.	Uses real-time decoder literature as the measurement standard, not as evidence for this implementation.	Device logs, bitstream/RTL hash and board vectors.

Table 33: Supported design trade-off and the evidence needed to strengthen each claim. The table summarizes what is established by the present software and analytical evidence, how that differs from adjacent GKP, learned-decoder and real-time-decoder work, and what would be required before making stronger hardware or channel-level claims.

Comparison axis	Representative prior-work metrics	Relevance to this paper
Analog GKP information	Analog syndrome decoding has been reported to lower required squeezing from about 16.0 dB to 9.8 dB in surface-GKP settings [3, 4].	The present work does not claim novelty in analog information. It uses recent analog histograms to update a low-dimensional affine correction surface.
Calibration-aware QEC	Prior calibration and drift-aware graph updates report gains ranging from percentage-level prior improvements to multi-fold mismatch reductions [13, 14, 15].	The calibrated object here is the physical-layer (K, b) surface consumed by the fast GKP correction path, not only an outer-code graph weight or decoder prior.
Logical-error and overhead targets	Surface-GKP studies report logical-failure and overhead targets such as 0.87% to 0.36% CNOT-failure reduction, 9.9 dB threshold behavior, and 291 modes plus 97 qubits versus 1457 bare qubits for a matched low-failure target [5].	The present metric is narrower: it is a protocol-defined q/p boundary-crossing proxy, not the same endpoint as prior-work LER or logical-failure rates. The result should therefore be read as a front-end calibration result, not as an end-to-end surface-GKP overhead claim.
Logical-channel fidelity and infidelity	Approximate-GKP analyses report logical-channel probabilities, loss-channel infidelity and near-optimal loss/amplification performance as channel-level figures of merit [31, 32, 33].	The current manuscript reports a residual-boundary Pauli-channel-style surrogate derived from q/p half-lattice crossings, but not finite-energy logical-channel fidelity or process infidelity. Fidelity-level claims would require channel-specific simulation, logical-channel probability estimates or process-infidelity reconstruction under a specified physical channel.

Table 34: External GKP and calibration comparison axes. The cited values are representative reported metrics from adjacent work, not normalized baselines for this study. They position the contribution and the missing validation without treating incompatible metrics as results of this manuscript.

Comparison axis	Representative prior-work metrics	Relevance to this paper
Real-device learned decoding	Learned decoders have reported real-processor LER reductions using large experimental and synthetic data sets [16, 18].	The present manuscript does not report real-processor learned-decoder LER. It uses software drift scenarios to test whether a learned slow loop can populate a deterministic affine surface.
Structured neural and soft decoders	Bosonic/GKP-like neural decoders, neural matching-weight predictors and hardware soft-information decoders provide related routes to improved LER or thresholds [19, 20, 21].	The present work does not claim novelty in neural decoding or soft information. Its distinction is that the neural component is confined to slow calibration of the affine physical correction surface.
Runtime pre-decoders and calibration-conditioned latency	AI pre-decoders and calibration-conditioned decoders report runtime anchors such as $38.78\ \mu\text{s}$ sub-step decoding or $85\text{--}95\ \mu\text{s}$ folded FiLM latency [17, 18].	The current latency argument is only analytical and software-emulated. A future device study must measure the affine fast path rather than borrow these per-shot neural latency values.
Real-time FPGA decoders	Hardware decoders report deployment metrics such as $28\text{--}42\ \text{ns}$ lookup latency and compressed lookup memory [24], $11.5\ \text{ns}$ per measurement round [25], sub- μs round or feedback timing [28, 29, 30], $550\ \text{ns}$ closed-loop latency and QPU-feedback integration [23], and scale, memory, area and power figures [26].	These works define the standard for hardware claims. This manuscript only specifies the required hardware measurement surface until the same classes of timing, resource and source-vs-board data are available.

Table 35: External learned-decoder and hardware-runtime comparison axes. The cited values are representative reported metrics from adjacent work, not normalized baselines for this study. They position the slow-loop calibration role and the missing hardware validation, not as results of this manuscript.

Material class	Material used in this manuscript	Interpretation limit
Predeclared simulation results	Four-scenario, five-mode ranking used in Table 12.	Supports relative ranking only inside the predeclared simulation protocol.
Ablation and mechanism materials	Feature and cross-seed intervention tables used in Tables 29 and 30.	Descriptive mechanism evidence; not a causal proof.
Statistical-calibration extension	Supplement-side comparison in Table 31.	Supports the affine-calibration contract; does not promote statistical calibration as a validated main comparator.
Controlled nearest-syndrome, oracle and wrapped-Gaussian checks	One-step and short-sequence local-Gaussian comparisons in Tables 17 and 22.	Supports estimator, hard-correction and decoder-limit sanity checks only; does not replace formal stronger baselines or <code>final_ler_mean</code> validation.
Holdout drift stress tests	Random-walk, burst/reset and faster-than-window controlled stress families in Table 23.	Supports non-hardware unseen-drift diagnostics only; does not establish trained-branch holdout generalization, confidence intervals or hardware robustness.
Residual-boundary channel surrogate	q/p half-lattice crossing event decomposition in Table 18 and method-level lattice sanity summary in Table 20.	Supports a Pauli-event surrogate bridge between residual-boundary scoring and channel language only; does not estimate finite-energy logical-channel fidelity, process fidelity or hardware fidelity.
Fast-path cost model	Analytical operation-count comparison in Table 26.	Supports low-complexity motivation only; does not establish FPGA timing, power or resource use.
Fixed-point parity check	Q4.20 software-emulation comparison in Table 27.	Supports numerical fixed-point feasibility under controlled samples only; does not establish source-vs-board agreement, synthesis, timing closure, resource use or power.
Literature metric crosswalk	Cited metric anchors for the external comparison table, including paper-level figure, table and equation anchors where available.	Supports related-work positioning only; not a normalized leaderboard and not experimental data generated by this study.
Training and runtime support	A clean CPU-only rerun and selected isolated <code>.tfLite</code> runtime checks.	Does not establish full reproducibility, default-environment portability or validated deployed operation.
Hardware validation requirements	Hardware measurement checklist without board measurements.	Does not establish board execution, timing, resource or source-vs-board agreement.

Table 36: Availability of supporting materials. Each row states how a material class is used in the manuscript and what stronger conclusion it cannot support.

Evidence layer	Supported statement	Not established by the present study
Predeclared simulation benchmark	Four-scenario, five-mode ranking under the predeclared protocol.	Expanded benchmark, <code>.tfLite</code> ranking or board ranking.
Ablation/mechanism	Descriptive feature sensitivity and cross-seed mechanism material.	Causal mechanism proof or intervention success.
Training/material	Canonical material chain and one clean CPU-only rerun.	Full reproducibility, cross-host portability or GPU/Linux reproducibility.
Isolated runtime	Selected preserved <code>.tfLite</code> artifacts execute in an isolated software runtime.	Default-environment compatibility, closed-loop hardware-in-the-loop validation or validated deployed operation.
Hardware validation	No board-level measurement is reported in this study.	Real-board execution success or hardware validation.

Table 37: Scope of reported manuscript claims.

Result class	Supporting table	Protocol anchor	Manuscript interpretation
Main simulation benchmark	Machine-readable comparison table and run summary.	Predeclared four-scenario, five-mode, paired-seed, repeats = 2 protocol.	Supports the predeclared ranking only, not an expanded benchmark, <code>.tflite</code> , real-board or paper-grade SOTA comparison.
Feature and teacher ablation	Ablation summary table and underlying data file.	Same simulation setting, with controlled feature or teacher inputs removed.	Supports feature-sensitivity wording, not teacher-parameter necessity or a proven causal mechanism.
Multi-seed mechanism summary	Six-seed intervention summary and underlying figure data file.	Matched descriptive intervention comparison across the listed seeds.	Supports descriptive mechanism evidence, not intervention success or a closed causal mechanism.
Supplementary statistical calibration	Statistical-calibration comparison rows and underlying data file.	Separate supplementary estimator protocol.	Supports the supplement-side statement that the affine contract admits a non-neural estimator; it is not part of the main ranked table.
Hardware validation requirements	No board timing/resource rows are reported; selected <code>.tflite</code> checks are used only as software-runtime support.	Future board validation must provide device logs, bitstream or firmware hashes, DMA evidence and source-vs-board vectors.	Defines future measurement requirements; it does not report board execution, validated deployed operation or source-vs-board agreement.

Table 38: Benchmark protocol and supporting-table map. The table records how the present figures and supplementary files are supported.

Theoretical object	Implementation bridge	Manuscript role	Missing before stronger claim
Syndrome stream s_t	Simulation sessions generate recent syndrome statistics from configured synthetic drift trajectories.	Defines the fast/slow-loop interface and the histogram-window inputs.	Physics-level syndrome wrapping and logical-channel validation.
Affine fast path	Runtime contract applies a clipped and quantized $\Delta_t = K_t s_t + b_t$.	Core method claim: complex estimation is kept outside the fast path.	Bit-accurate fixed-point and board-latency evidence.
Effective drift state	Four predeclared synthetic drift families instantiate static, ramp, step and periodic perturbations of $(\sigma, \mu_q, \mu_p, \vartheta)$.	Controlled synthetic probes of adaptation demands.	Machine-readable field-to-symbol mapping and holdout drift families.
Main metric	The comparison table reports final_ler_mean means and descriptive standard deviations for paired seeds with repeats = 2.	Predeclared simulation ranking.	Confidence intervals, larger repeat count and stronger logical-error validation.
Hardware target	Board-facing entries specify future measurement requirements.	Future board-validation protocol only.	Device logs, bitstream or RTL hashes, DMA evidence and timing/resource data.

Table 39: Theory-to-configuration bridge for the manuscript. The bridge records how the notation in the method sections is instantiated by the predeclared simulation protocol, while making clear that the present data are not a physical-noise or hardware-validation result.

Layer	Artifact or runtime class	Manuscript use	Not established by the present study
Main comparison	controlled simulation rows from <code>comparison.csv</code> and <code>summary.json</code> .	Predeclared four-scenario, five-mode ranking.	True <code>.tflite</code> , real-board or expanded benchmark ranking.
Hybrid residual branch	Preserved <code>.npz</code> model artifact recorded in the hybrid rows.	Slow-loop model-artifact evidence for the current hybrid controlled simulation branch.	Portable deployment or board execution.
Baseline modes	controlled simulation modes with no model artifact path in the rows.	Classical and residual baselines under the same protocol.	Learned artifact parity or <code>.tflite</code> baseline evidence.
Software runtime check	Selected preserved <code>.tflite</code> files executed in an isolated software runtime; JSON metadata artifacts are not treated as executables.	Supporting runtime evidence only.	Default-environment compatibility, closed-loop hardware-in-the-loop validation or validated deployed operation.
Non-executable metadata artifacts	JSON-sidecar or metadata-only artifacts in the runtime abstraction.	Artifact taxonomy only.	True TFLite execution.
Hardware validation requirements	Validation requirements without a current board measurement.	Future hardware measurement requirements.	Real-board execution success or hardware validation.
Statistical calibration	Supplementary statistical-calibration supporting table.	Evidence that the affine contract can host a non-neural estimator.	A validated main-comparator role or replacement of the main table.

Table 40: Runtime-artifact validation scope. The table separates simulation results, model artifacts, true `.tflite` support, non-executable metadata artifacts, hardware-validation requirements and supplementary statistical-calibration evidence so that the manuscript does not upgrade one artifact class into another.

Coverage group	Checked manuscript material	Interpretation limit
Main performance tables	Main results, paired deltas, paired-uncertainty and LER advantage-margin underlying rows are checked against underlying data tables.	Descriptive two-repeat software comparison only; not confidence intervals, standard errors, p-values, significance tests or an expanded benchmark.
Controlled diagnostics	Oracle-affine, logical-surrogate, surrogate-fidelity, lattice-channel sanity, finite-energy toy-channel, boundary-sensitivity, sequence-baseline, holdout-stress, affine local-validity and commit-lag tables are checked against generated CSV files.	Diagnostic layers only; not trained-branch holdout proof, calibrated finite-energy logical-channel fidelity or hardware latency.
Implementation feasibility	Fast-path cost, fixed-point parity, runtime discipline and validation-contract rows are checked against generated analysis tables or figure data tables.	Analytical/software feasibility only; not FPGA synthesis, timing, resource, power or source-vs-board evidence.
Source and literature maps	Metric-readiness, literature crosswalk, closest-work positioning, row metadata, figure manifest and file-level supporting manifest are checked for row, citation or hash consistency.	Not a leaderboard, not a full independent reproduction package.
Completed paired-interval scenarios	The completed <code>static.bias.theta</code> and <code>linear.ramp</code> formal paired-interval rows are generated from the repeat summary and checked against the manuscript interval table.	Two completed scenarios only; not all-scenario repeat-expanded evidence, <i>p</i> -value evidence, holdout robustness or hardware validation.
Validation-protocol material	Benchmark-expansion, runner-smoke, availability, validation-scope, runtime-artifact and hardware-plan material are classified separately.	Protocol or requirements material only; not result evidence or hardware validation.

Table 41: Supporting-table coverage matrix. The matrix identifies which manuscript table families are backed by checked analysis files and which are protocol or requirements material. It is a traceability aid, not an additional result.

Layer	Evidence available	Manuscript use	Missing for stronger claim
Main comparison	Predeclared four-scenario, five-mode controlled simulation benchmark with paired seeds, repeats = 2, paired-delta and resampling-envelope tables.	Predeclared-set ranking; hybrid residual branch has the lowest mean <code>final_ler_mean</code> in all four scenarios; paired tables provide data transparency only.	Paired CI, more repeats and a formal expanded holdout benchmark.
Error bars	Figure supporting data contain mean plus descriptive SD for Fig. 2.	Visual uncertainty marker only.	Do not treat as CI, standard error or statistical significance.
Supporting files	<code>comparison.csv</code> , <code>summary.json</code> , figure underlying data tables, figure manifest, consistency report and file-level supporting-file manifest plus row metadata for the main controlled simulation rows.	Table and figure support for the manuscript.	Whole-manuscript supporting-file check and independent reproduction.
Reproducibility	Recorded run metadata, host checks and scoped training/material support.	Available repository-artifact reproducibility.	Repeated runs, cross-host checks or containerized reproduction.
Hardware path	No current board measurement is available.	Future measurement surface only.	Device logs, bit-stream/RTL hash, DMA evidence, latency and resource data, and source-to-board vectors.

Table 42: Statistical treatment and reproducibility. The current data support a scoped controlled simulation ranking with descriptive uncertainty; stronger statistical, reproducibility or hardware claims require the missing evidence listed in the rightmost column.

Evidence class	Present evidence	Supported interpretation	Evidence needed for stronger claims
Descriptive benchmark	Four-scenario, five-mode paired-seed comparison.	Hybrid- <i>b</i> has the lowest descriptive mean final_1er_mean in the predeclared controlled simulation rows.	The data support a descriptive ranking only; statistical significance, broad robustness, hardware latency and deployment readiness require additional evidence.
Runner and row-accounting check	Shortened UKF-versus-Hybrid rows complete for the four predeclared scenarios.	The repeat-expansion command path and row accounting were exercised.	The shortened rows are an execution-path check, not manuscript performance evidence or inferential uncertainty.
Completed paired-interval rows	Two formal-length scenario rows have completed with 12 unique paired rows and positive paired-interval lower bounds.	The completed scenarios support scenario-level positive paired-interval statements; all-scenario and pooled wording require additional formal rows.	Complete all four scenarios with at least 12 and preferably 16 paired repeats, no missing paired rows and positive predeclared paired-interval lower bounds per scenario plus pooled analysis.
Holdout drift validation	Specified for future study.	Holdout drift families are defined as validation families.	Run the predeclared random-walk, burst/reset and faster-than-window families with missing-run accounting; do not treat controlled stress diagnostics as holdout generalization proof.
Hardware-facing measurements	Outside the present scope.	Hardware latency, resource, power and source-vs-board agreement remain measurement targets.	Provide board logs, bitstream or RTL hash, source vectors, measured latency/resource/power and source-vs-board agreement before claiming hardware timing closure or board-level correction success.

Table 43: Validation thresholds for stronger statistical and hardware claims. The shortened runner check and incomplete holdout route are not performance evidence. Stronger wording requires the evidence listed in the final column.

Scenario	Paired repeats	Mean delta	Paired- <i>t</i> 95% interval	Bootstrap 95% interval
static_bias_theta	12	0.015563	[0.014778, 0.016348]	[0.014933, 0.016256]
linear_ramp	12	0.022417	[0.021215, 0.023619]	[0.021405, 0.023440]

Table 44: Completed scenario-level paired-interval checks. The rows use the completed formal-length **static_bias_theta** and **linear_ramp** UKF-minus-Hybrid **final_1er_mean** deltas. They support positive interval statements for these scenarios, not all-scenario repeat-expanded advantage, holdout robustness, *p*-value evidence or hardware validation.

Manuscript claim	Supporting material	Reporting treatment
Drift adaptation can be expressed as a runtime affine calibration problem.	Problem formulation, architecture contract and predeclared controlled simulation protocol.	Treated as a method claim; supporting material consists of definitions, protocol metadata and parameter-flow schematic, not performance measurements.
The main-table hybrid residual branch has the lowest mean <code>final_ler_mean</code> in the predeclared four-scenario set.	Table 12 from the predeclared controlled simulation benchmark.	Machine-readable per-scenario rows are retained; seed-level uncertainty would be needed before claiming statistical robustness.
Feature choices affect residual calibration behavior.	Table 29; scoped feature and teacher ablations.	Treat as an appendix supporting table, not as a teacher-necessity claim.
The mechanism story remains descriptive.	Table 30; six-seed intervention summary with mixed or harmful intervention outcomes.	Link the figure data files and trace summaries; the result is not a causal mechanism identification.

Table 45: Support map for the primary manuscript claims. The table records how method, main-result and mechanism statements are tied to supporting materials.

Manuscript claim	Supporting material	Reporting treatment
The metric set is intentionally layered.	Table 3; <code>final_1er_mean</code> , controlled checks, the residual-boundary Pauli-event surrogate, cost counts and fixed-point software parity are reported metrics, while finite-energy channel fidelity and hardware timing remain external standards.	The supporting CSV is retained with the manuscript; <code>final_1er_mean</code> and the surrogate average-fidelity field are not relabelled as channel fidelity or measured latency.
Controlled nearest-syndrome, oracle-affine and wrapped-Gaussian checks are diagnostic baselines.	Tables 17 and 22, plus the derived local-validity diagnostic in Table 24; controlled local-Gaussian underlying data tables.	The one-step and short-sequence protocols remain separate from the predeclared controlled simulation ranking; they are not a full nearest-lattice, known-noise or wrapped-decoder benchmark.
Holdout drift stress tests probe adaptation space, not trained-branch generalization.	Table 23; controlled random-walk, burst/reset and faster-than-window diagnostic data table.	These rows are non-hardware stress diagnostics, not expanded benchmark evidence, confidence intervals or robustness proof.
Runtime counters and the validation-contract figure support software observability.	Table 28 and Fig. 3; runtime-discipline and validation-contract data tables.	These rows support software protocol observability, not board latency, hardware reliability, synthesis, timing closure or source-vs-board agreement.
The fixed-point claim is numerical, not hardware-complete.	Table 27; software-emulation data table; Q4.20 emulation introduces no residual-boundary crossing change in the controlled sample set.	This is a numerical-emulation result, not source-vs-board agreement, timing closure, resource measurement, power measurement or real-board latency.
The affine contract can host non-neural calibration.	Table 31; supplementary statistical-calibration analysis.	The supporting rows are supplement-side and do not replace the main comparator table or identify a unique threshold.
Hardware validation requires direct device data.	Table 32 and the absence of board measurements in this study.	Hardware claims require board logs, bitstream hashes, latency/resource data and source-vs-board vectors.

Table 46: Support map for diagnostic, runtime and validation statements. The table continues the reporting route for controlled diagnostics, numerical checks, supplementary calibration and hardware-measurement requirements.

Validation concern	Present manuscript treatment	Evidence still needed for a stronger claim
The benchmark is too narrow.	The claim is restricted to the predeclared four-scenario, five-mode controlled simulation protocol, and the protocol is reported as a predeclared simulation benchmark.	A predeclared expanded benchmark with holdout drift families, stronger baselines and more repeats or uncertainty estimates.
The method lacks strong theoretical baselines.	The text adds controlled one-step and short-sequence nearest-syndrome / oracle-affine / wrapped-Gaussian checks and does not claim superiority over nearest-lattice or fully tuned sequence-level wrapped decoders.	A formal known-noise affine, nearest-lattice or wrapped-Gaussian baseline under a predeclared benchmark protocol.
The mechanism is not causal.	Ablation and intervention materials are described as feature sensitivity and descriptive mechanism evidence only.	Mechanism-specific perturbations, trace-level causal isolation or a predictive failure-mode test.
Hardware claims require direct measurements.	Board-level validation is stated as a future measurement surface because no board measurements are reported here.	Device logs, bitstream hashes, fixed-point agreement, latency/resource measurements and source-vs-board tests.
The neural branch is not the only plausible estimator.	Statistical calibration is kept as a supplementary analysis that supports the affine contract rather than a main comparator.	A future predeclared comparator selection protocol if the statistical estimator is to be used in the main ranking.
The work may not be reproducible outside the present software environment.	Reproducibility statements are restricted to available analysis records and single-host checks.	Repeated-run, cross-host or containerized reproduction with explicit dependency locks and artifact manifests.

Table 47: Limitations and required validation for stronger claims. Each row states how the manuscript limits the present claim and what evidence would be required before stronger wording is justified.

Term	Meaning in this manuscript	Boundary
Runtime-consistent affine calibration	A slow loop updates (K_t, b_t) , while the fast path applies $\Delta_t = K_t s_t + b_t$.	Full posterior decoding or hardware validation.
Fast path	The per-shot deterministic affine correction surface.	A measured FPGA implementation, unless future board evidence is added.
Slow calibration loop	A teacher, neural residual branch or statistical estimator that updates affine parameters from recent syndrome statistics.	A requirement that the slow loop must be neural.
Teacher-anchored hybrid residual branch	The predeclared main-table hybrid mode; ablation rows do not establish teacher-parameter necessity.	A universal SOTA decoder or causal mechanism proof.
Supplementary statistical calibration	A supplement-side check that non-neural estimators can also satisfy the affine contract.	A main-comparator role.
controlled simulation	A source-controlled software protocol for testing the runtime calibration interface under synthetic drift.	Real-board HIL or deployment success.
Hardware validation check	A statement that no board measurement is reported and that future measurements must provide timing, resource and source-vs-board data.	Successful board execution.

Table 48: Reader-facing terminology and claim boundaries. The table defines how recurrent terms are used in the manuscript.

Figure	Main conclusion	Required supporting files	Reporting basis
Fig. 1 architecture	GKP drift adaptation is expressed as a dual-loop affine runtime contract.	Method schematic, parameter-flow definition and validation scope notes.	Python manuscript figure with schematic support map.
Fig. 2 main controlled simulation results	Hybrid residual calibration is best in the predeclared controlled simulation table.	Main comparison data table, final_ler_mean mean, descriptive SD and later paired comparison.	Python manuscript figure with mean + descriptive SD supporting data; CI still missing.
Fig. 3 ablation and mechanism	Feature and histogram dynamics constrain the mechanism story.	Ablation rows, cross-seed intervention rows and trace summaries.	Python manuscript figure with descriptive-only supporting data.
Fig. 4 statistical calibration	Non-neural calibration can occupy the same affine contract without becoming the main comparator.	Supplementary statistical-calibration rows and no-main-comparator notes.	Python manuscript figure; supplementary analysis only.
Fig. 5 validation contract	The present validation package supports a software fast-path contract across analytical arithmetic, Q4.20 parity and runtime counter observability.	Validation-contract data table plus fast-path cost, fixed-point parity and runtime-discipline data tables.	Manuscript figure with data-backed software summary only; no FPGA timing, resource, source-vs-board or board-latency measurement.

Table 49: Figure data and reporting basis. Each figure is tied to the source files or validation record needed to support its intended conclusion.